


# 昇腾950 NPU架构白皮书深度解读

主讲：李艳华、顾宝成、刘锡明

# 昇腾950的架构设计基于哪些重要思考？



# 昇腾950PR/DT 芯片规格




| 规格项    |                           | 昇腾950PR                      | 昇腾950DT                          |
|--------|---------------------------|------------------------------|----------------------------------|
| 算力     | Cube算力规格                  | 378T<br>(FP16/BF16)          | 432~486T<br>(FP16/BF16)          |
|        |                           | 757T<br>(MxFP8, HiF8, MxFP6) | 864~972T<br>(MxFP8, HiF8, MxFP6) |
|        |                           | 1514T<br>(MxFP4)             | 1728~1944T<br>(MxFP4)            |
|        | Vector算力规格<br>(FP16/BF16) | 47T                          | 54~60.8T                         |
| Memory | 总算力规格<br>(FP16/BF16)      | 425T                         | 486~547T                         |
|        | Memory容量<br>(GB)          | 128                          | 96/144                           |
|        | Memory带宽                  | 1.6TB/s                      | Up to 4TB/s                      |
| IO规格   | UB Link                   | Up to 2016GB/s               | 2016GB/s                         |
|        | PCIE                      | 128GB/s<br>RC & EP           | 128GB/s<br>RC & EP               |
|        | RoCE/UBoE                 | 400Gbps                      | 400Gbps                          |



# 昇腾950采用了第三代达芬奇架构，有哪些创新？

# 昇腾950采用了第三代达芬奇架构



## Cube MMAD计算支持的数据类型:


1. FP16/BF16/HF32/FP32/S8
2. FP8(E5M2/E4M3)
3. MXFP8(E5M2/E4M3)
4. MXFP4(E2M1/E1M2)
5. HiF8

## 950相比910系列数据通道主要改动点:

1. 增加LOC -> UBuffer数据通路
2. 增加UBuffer -> L1数据通路
3. 增加SSBuf支持CV之间的消息通路

# 昇腾950如何突破“内存墙”？


# 昇腾950 Memory内存层次





# 昇腾950 Memory层次中主要Memory及其大小


| Memory              | Memory size         |
|---------------------|---------------------|
| L1 Buffer           | 512KB per AI Core   |
| LOA Buffer          | 64KB per AI Core    |
| LOB Buffer          | 64KB per AI Core    |
| LOC Buffer          | 256KB per AI Core   |
| Unified Buffer (UB) | 512KB per AI Core   |
| CPU L1 Cache        | 64KB per CPU Core   |
| CPU L2 Cache        | 1MB per CPU Core    |
| L3 Cache            | 4MB per CPU Cluster |
| L2 Cache            | Up to 128MB         |
| 昇腾950PR片上内存         | Up to 128GB         |
| 昇腾950DT片上内存         | Up to 96/144GB      |



L2 hint 典型应用场景示意图



# 从单芯片到超节点——互联、通信与集群调度

# Ascend950超大互联带宽




- 2 IO DIE
- 每IO DIE 9端口
- 每端口4 Lane
- 18个端口最大支持2016 GB/s带宽

# Ascend950 ScaleUp网络互联




# Ascend950 ScaleOut网络互联


UB\_RTP互联




UBOE互联



ROCE互联



# Ascend950 CCU加速



CCU 核心能力

## 扫码获取白皮书& 互动拿社区礼品



<https://gitcode.com/cann>

**CANN**



# Thank you.

社区愿景：打造开放易用、技术领先的AI算力新生态

社区使命：使能开发者基于CANN社区自主研究创新，构筑根深叶茂、跨产业协同共享共赢的CANN生态

Vision: Building an Open, Easy-to-Use, and Technology-leading AI Computing Ecosystem

Mission: Enable developers to independently research and innovate based on the CANN community and build a win-win CANN ecosystem with deep roots and cross-industry collaboration and sharing.



上CANN社区获取干货



关注CANN公众号获取资讯