

RoPE-Matrix

昇腾亲和的RoPE算法实现与算子开源

中央媒体技术院-图像工程部 陈敏琪
2025年12月20日



Content 目录

- 01 RoPE-Matrix算法原理
- 02 融合算子开发与贡献

RoPE-Matrix算法原理

01

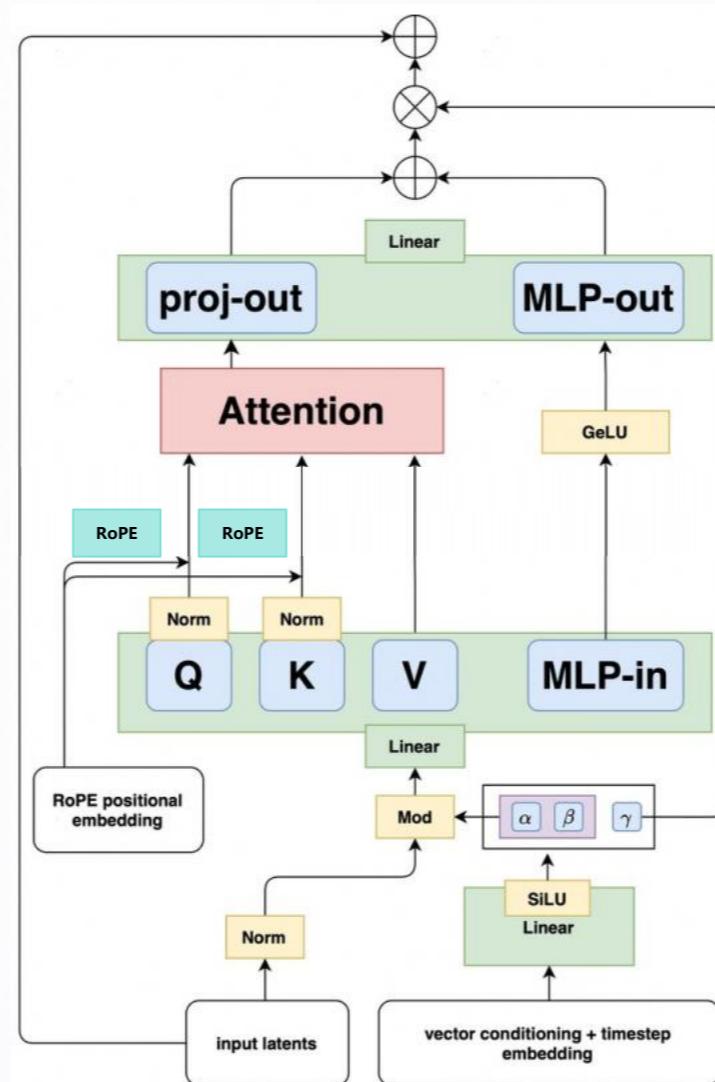
RoPE-Matrix：昇腾亲和的RoPE算法实现

多模态生成推理性能分析：Vector占比高，其中RoPE相关计算占到Vector的44%开销。

FLUX.1-Kontext性能分析

性能拆解	910B	占比
Cube类算子 (MatMul等)	0.452	42.6%
CV融合算子 (FlashAttention)	0.348	34.8%
Vector算子	0.238	22.5%
端到端耗时 【秒】	1.059	100%

A fused DiT block equipped with rotary positional embeddings



基于单个block的vector前向计算打开分析：**RoPE占到近一半开销**

类别	耗时[ms]	占比
RoPE	0.110	44%
GeLU	0.026	11%
LN	0.022	9%
TransData	0.021	8%
Concat	0.021	8%
Mul	0.013	6%

RoPE-Matrix: 昇腾亲和的RoPE算法实现

2021.04: Roformer-Enhanced Transformer With Rotary Position Embedding 提出当前广泛使用的RoPE，旋转位置编码，**应考虑任意两个token直接的相对位置信息**，被广泛应用于LLaMA、ChatGLM、PaLM等当今最先进的大模型中。

①

$$f_q(\mathbf{x}_m, m) = (\mathbf{W}_q \mathbf{x}_m) e^{im\theta}$$

$$f_k(\mathbf{x}_n, n) = (\mathbf{W}_k \mathbf{x}_n) e^{in\theta}$$

$$g(\mathbf{x}_m, \mathbf{x}_n, m - n) = \text{Re}[(\mathbf{W}_q \mathbf{x}_m)(\mathbf{W}_k \mathbf{x}_n)^* e^{i(m-n)\theta}]$$

考虑token的dim=2场景，第m、n个token之间的相关性分数，可以通过**复平面上的旋转**，保证了**相对位置信息的显式编码及稳定性**（不改tensor的模长）。

②

$$f_{\{q,k\}}(\mathbf{x}_m, m) = \begin{pmatrix} \cos m\theta & -\sin m\theta \\ \sin m\theta & \cos m\theta \end{pmatrix} \begin{pmatrix} W_{\{q,k\}}^{(11)} & W_{\{q,k\}}^{(12)} \\ W_{\{q,k\}}^{(21)} & W_{\{q,k\}}^{(22)} \end{pmatrix} \begin{pmatrix} \mathbf{x}_m^{(1)} \\ \mathbf{x}_m^{(2)} \end{pmatrix}$$

复数计算对软硬件支持有要求，可以转换成**旋转矩阵乘**实现等价功能

③

$$f_{\{q,k\}}(\mathbf{x}_m, m) = \mathbf{R}_{\Theta, m}^d \mathbf{W}_{\{q,k\}} \mathbf{x}_m$$

$$\mathbf{R}_{\Theta, m}^d = \begin{pmatrix} \cos m\theta_1 & -\sin m\theta_1 & 0 & 0 & \dots & 0 & 0 \\ \sin m\theta_1 & \cos m\theta_1 & 0 & 0 & \dots & 0 & 0 \\ 0 & 0 & \cos m\theta_2 & -\sin m\theta_2 & \dots & 0 & 0 \\ 0 & 0 & \sin m\theta_2 & \cos m\theta_2 & \dots & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & 0 & 0 & \dots & \cos m\theta_{d/2} & -\sin m\theta_{d/2} \\ 0 & 0 & 0 & 0 & \dots & \sin m\theta_{d/2} & \cos m\theta_{d/2} \end{pmatrix}$$

对于 token 的 dim=d 场景，可以**对d做两两分组**，分别做不同角度的旋转，实现不同的相对位置编码

$$\Theta = \{\theta_i = 10000^{-2(i-1)/d}, i \in [1, 2, \dots, d/2]\}$$

④

$$\mathbf{R}_{\Theta, m}^d \mathbf{x} = \begin{pmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_{d-2} \\ x_{d-1} \end{pmatrix} \otimes \begin{pmatrix} \cos m\theta_0 \\ \cos m\theta_0 \\ \cos m\theta_1 \\ \cos m\theta_1 \\ \vdots \\ \cos m\theta_{d/2-1} \\ \cos m\theta_{d/2-1} \end{pmatrix} + \begin{pmatrix} -x_1 \\ x_0 \\ -x_3 \\ x_2 \\ \vdots \\ -x_{d-1} \\ x_{d-2} \end{pmatrix} \otimes \begin{pmatrix} \sin m\theta_0 \\ \sin m\theta_0 \\ \sin m\theta_1 \\ \sin m\theta_1 \\ \vdots \\ \sin m\theta_{d/2-1} \\ \sin m\theta_{d/2-1} \end{pmatrix}$$

由于R矩阵的**稀疏性**，直接用矩阵乘法实现会很浪费算力，可以转换为**逐元素Vector乘实现**

⑤

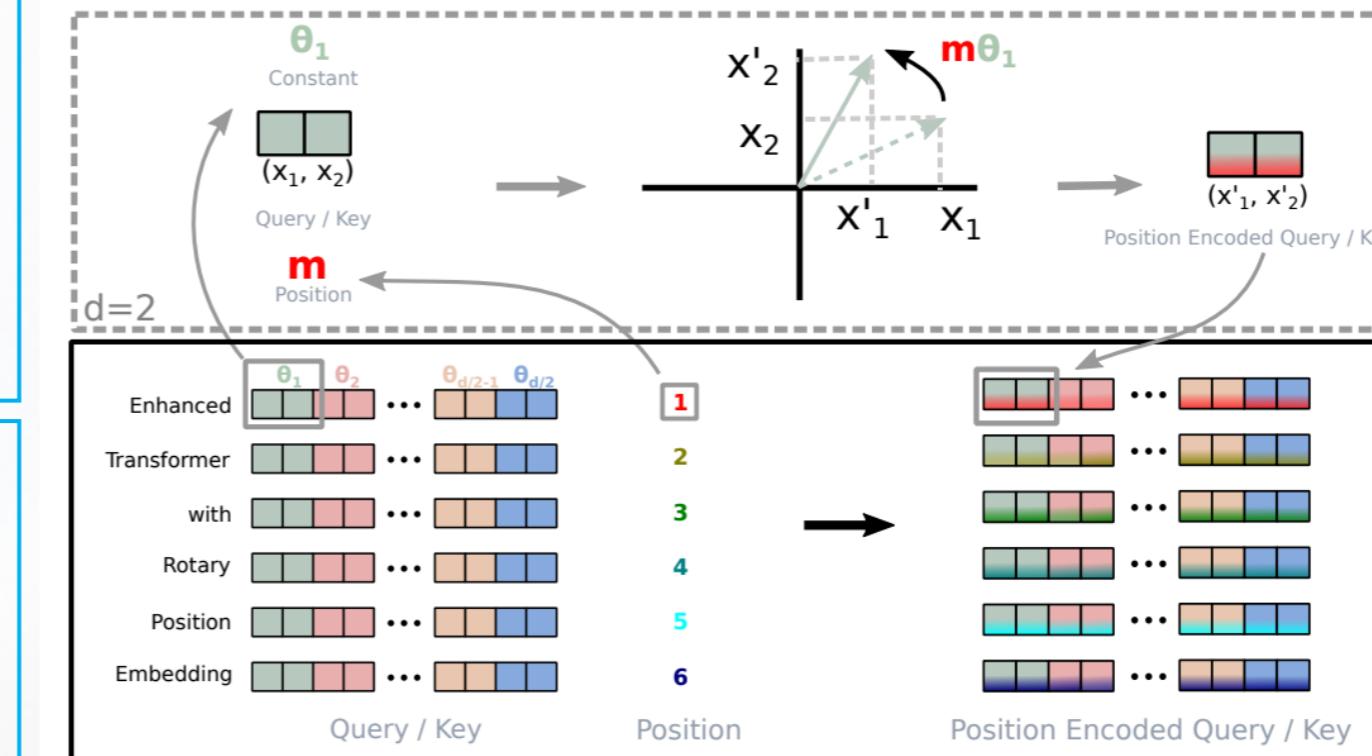
- half模式：

```
x1, x2 = torch.chunk(input, 2, -1)
x_new = torch.cat((-x2, x1), dim=-1)
output = r1 * input + r2 * x_new
```

根据公式④的代码实现。实际社区中会分为**half**及**interleave**两种模式，区别仅在于 dim 维度两两分组的策略不同，数学意义等价。

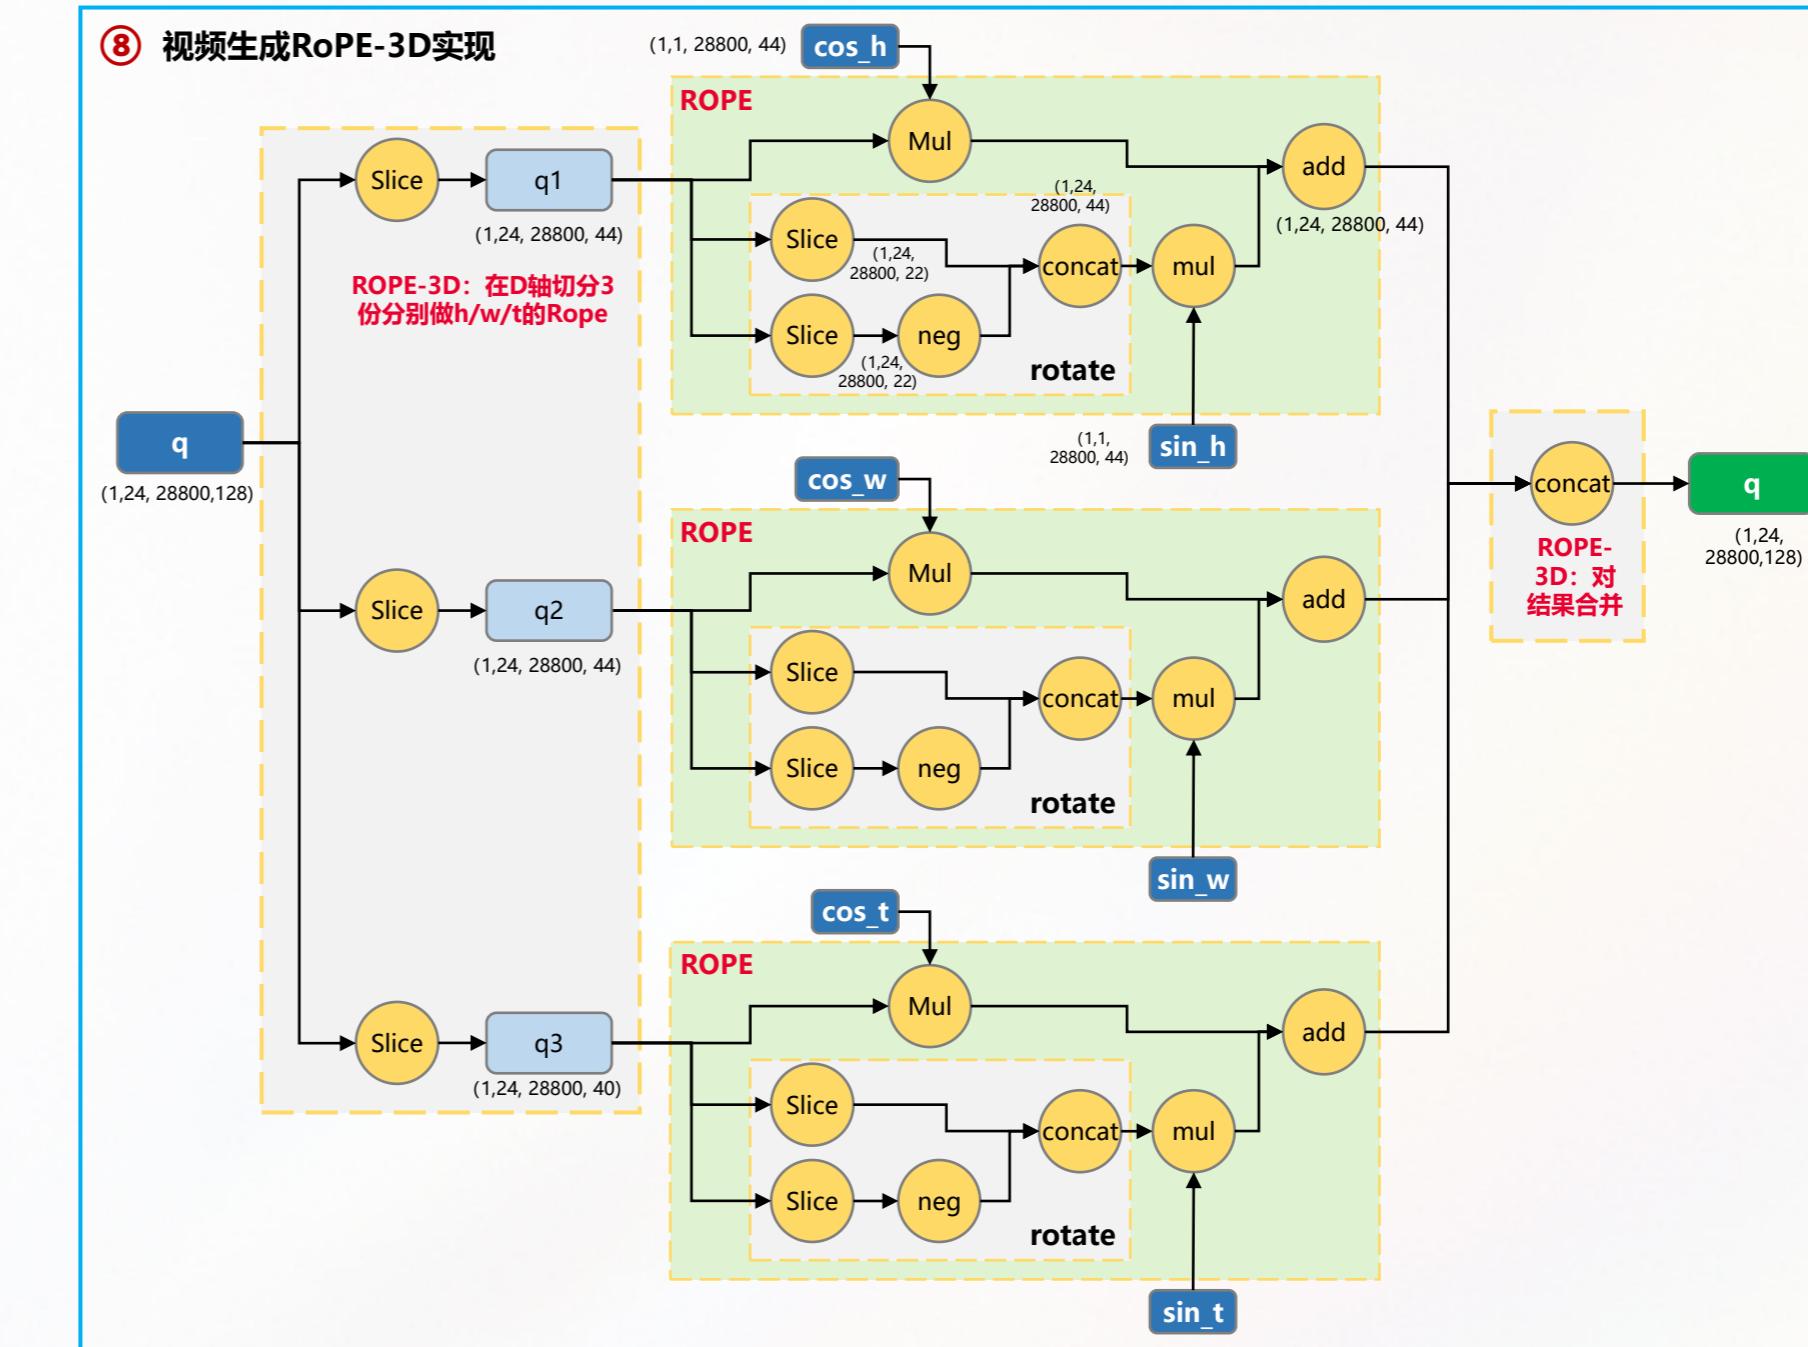
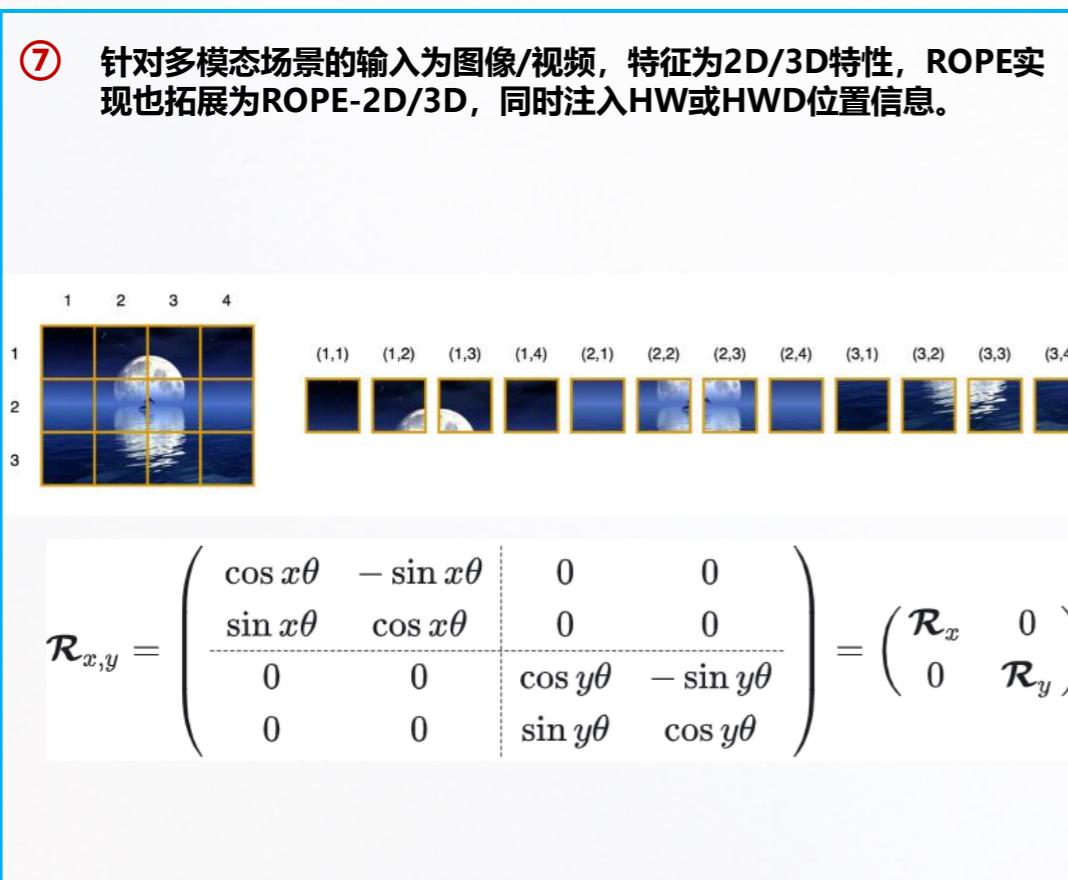
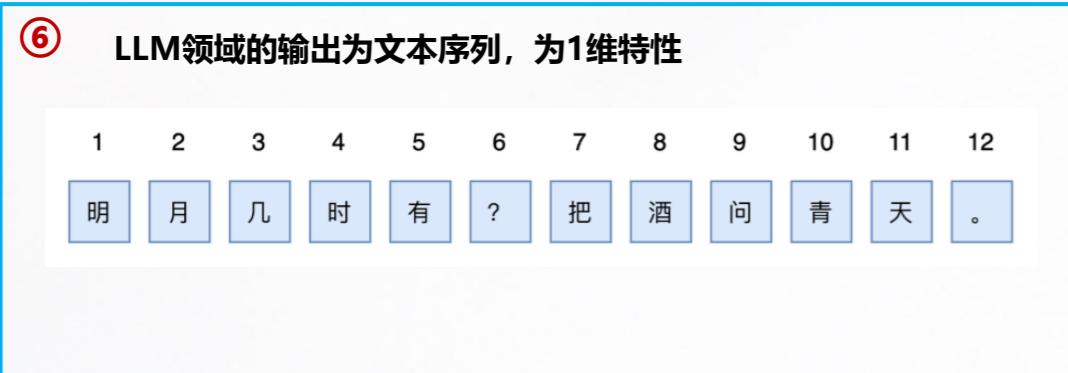
- interleave模式：

```
x1 = input[..., ::2]
x2 = input[..., 1::2]
x_new = rearrange(torch.stack((-x2, x1), dim=-1), "... d two -> ... (d two)", two=2)
output = r1 * input + r2 * x_new
```



RoPE-Matrix: 昇腾亲和的RoPE算法实现

2024.03: VisionLLaMA: A Unified LLaMA Backbone for Vision Tasks提出RoPE-2D, **位置编码信息考虑图像Patch的x,y坐标**, 当前被广泛应用于多模态大模型中。

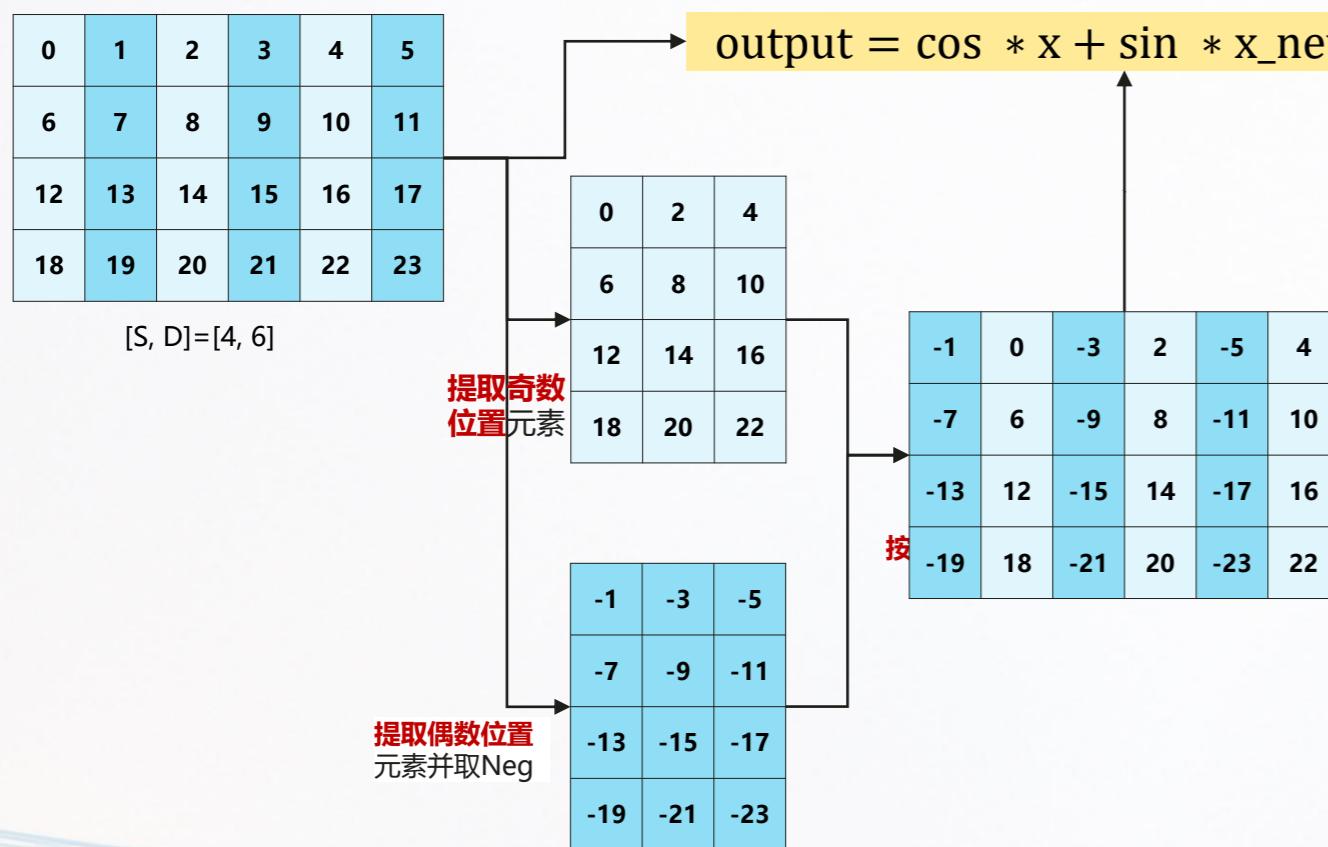


RoPE-Matrix: 昇腾亲和的RoPE算法实现

Interleave模式

问题1: interleave模式涉及到对元素奇偶数位操作, 昇腾不亲和;

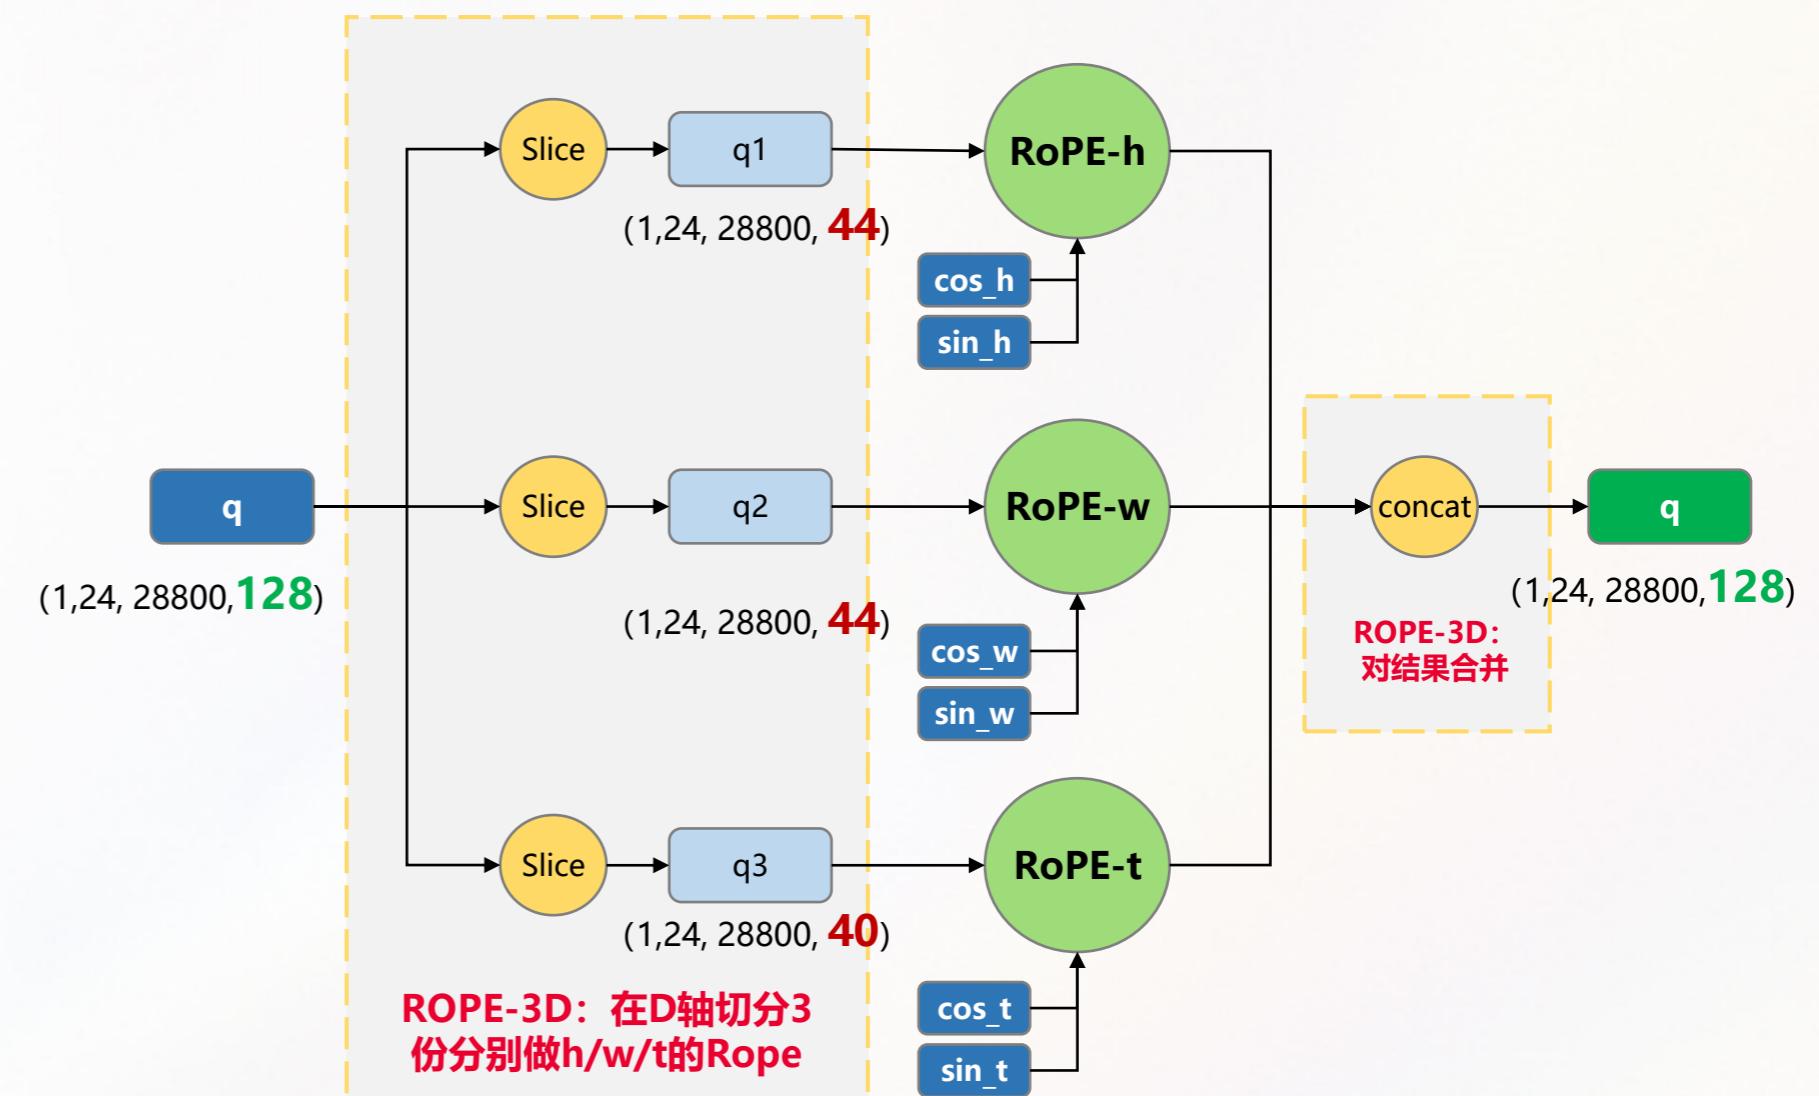
$$R_{\Theta, m}^d \mathbf{x} = \begin{pmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_{d-2} \\ x_{d-1} \end{pmatrix} \otimes \begin{pmatrix} \cos m\theta_0 \\ \cos m\theta_0 \\ \cos m\theta_1 \\ \cos m\theta_1 \\ \vdots \\ \cos m\theta_{d/2-1} \\ \cos m\theta_{d/2-1} \end{pmatrix} + \begin{pmatrix} -x_1 \\ x_0 \\ -x_3 \\ x_2 \\ \vdots \\ -x_{d-1} \\ x_{d-2} \end{pmatrix} \otimes \begin{pmatrix} \sin m\theta_0 \\ \sin m\theta_0 \\ \sin m\theta_1 \\ \sin m\theta_1 \\ \vdots \\ \sin m\theta_{d/2-1} \\ \sin m\theta_{d/2-1} \end{pmatrix}$$



RoPE-3D模式

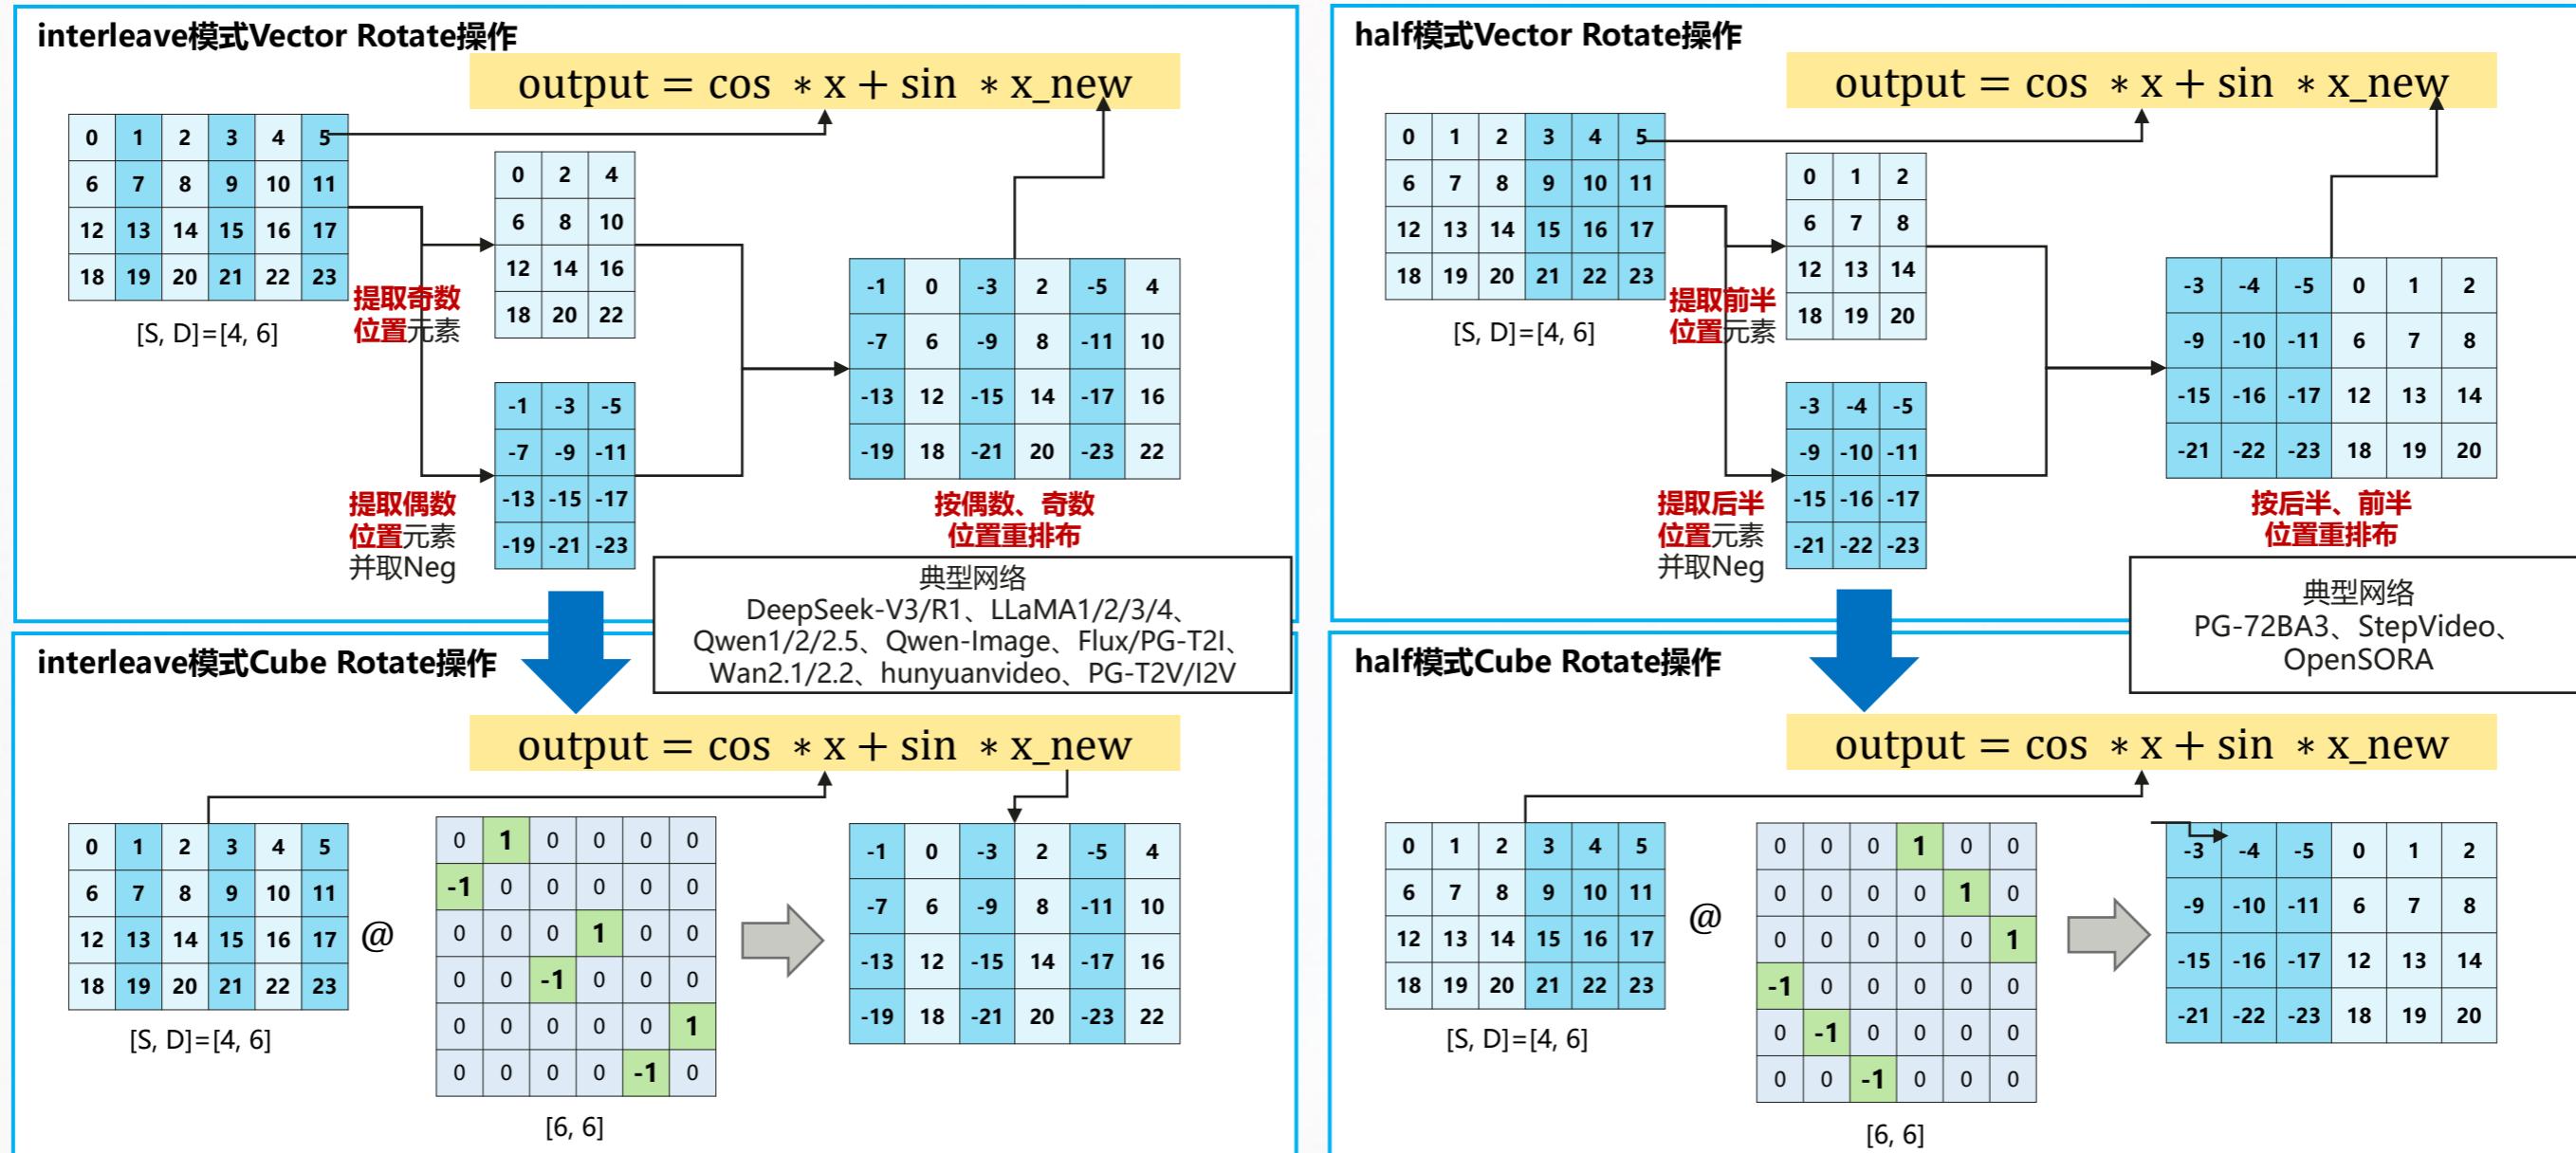
问题2: RoPE-3D模式下, 会对数据进行切片与Concat操作, 引入新开销

问题3: 切片后的数据, 尾轴为44/44/40, 后续RoPE-h/w/t计算/搬运效率低



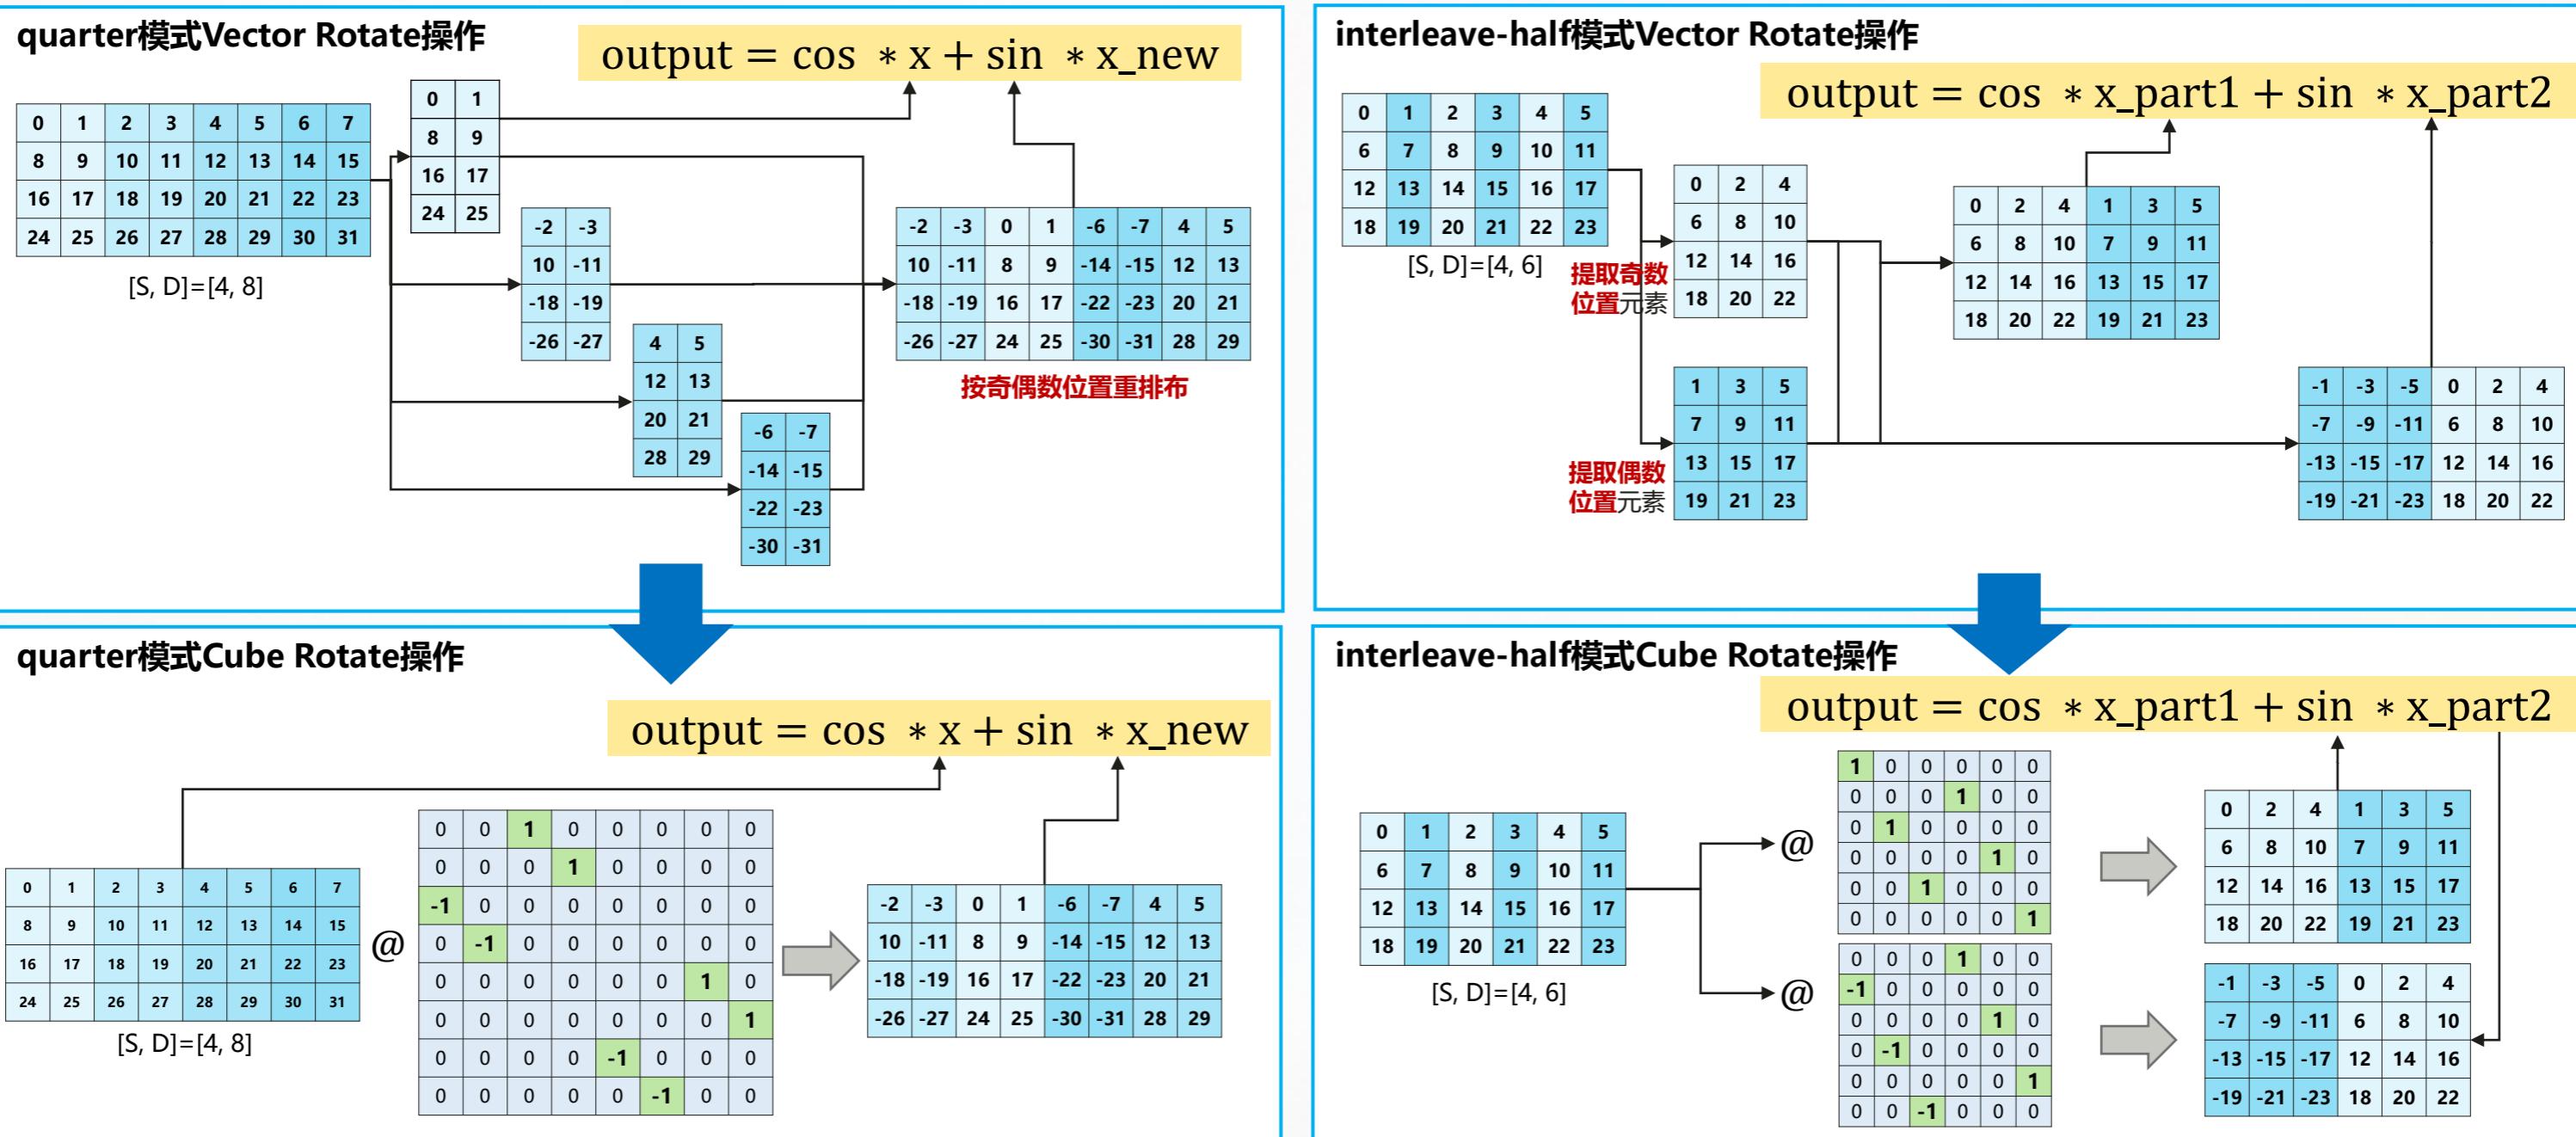
RoPE-Matrix: 昇腾亲和的RoPE算法实现

Step1：以Cube计算替代Vector计算的新RoPE实现：针对原RoPE中存在大量的对数据切分与合并操作，通过构造对应的操作矩阵以矩阵乘实现相关功能，以此替代昇腾不亲和的Vector操作。



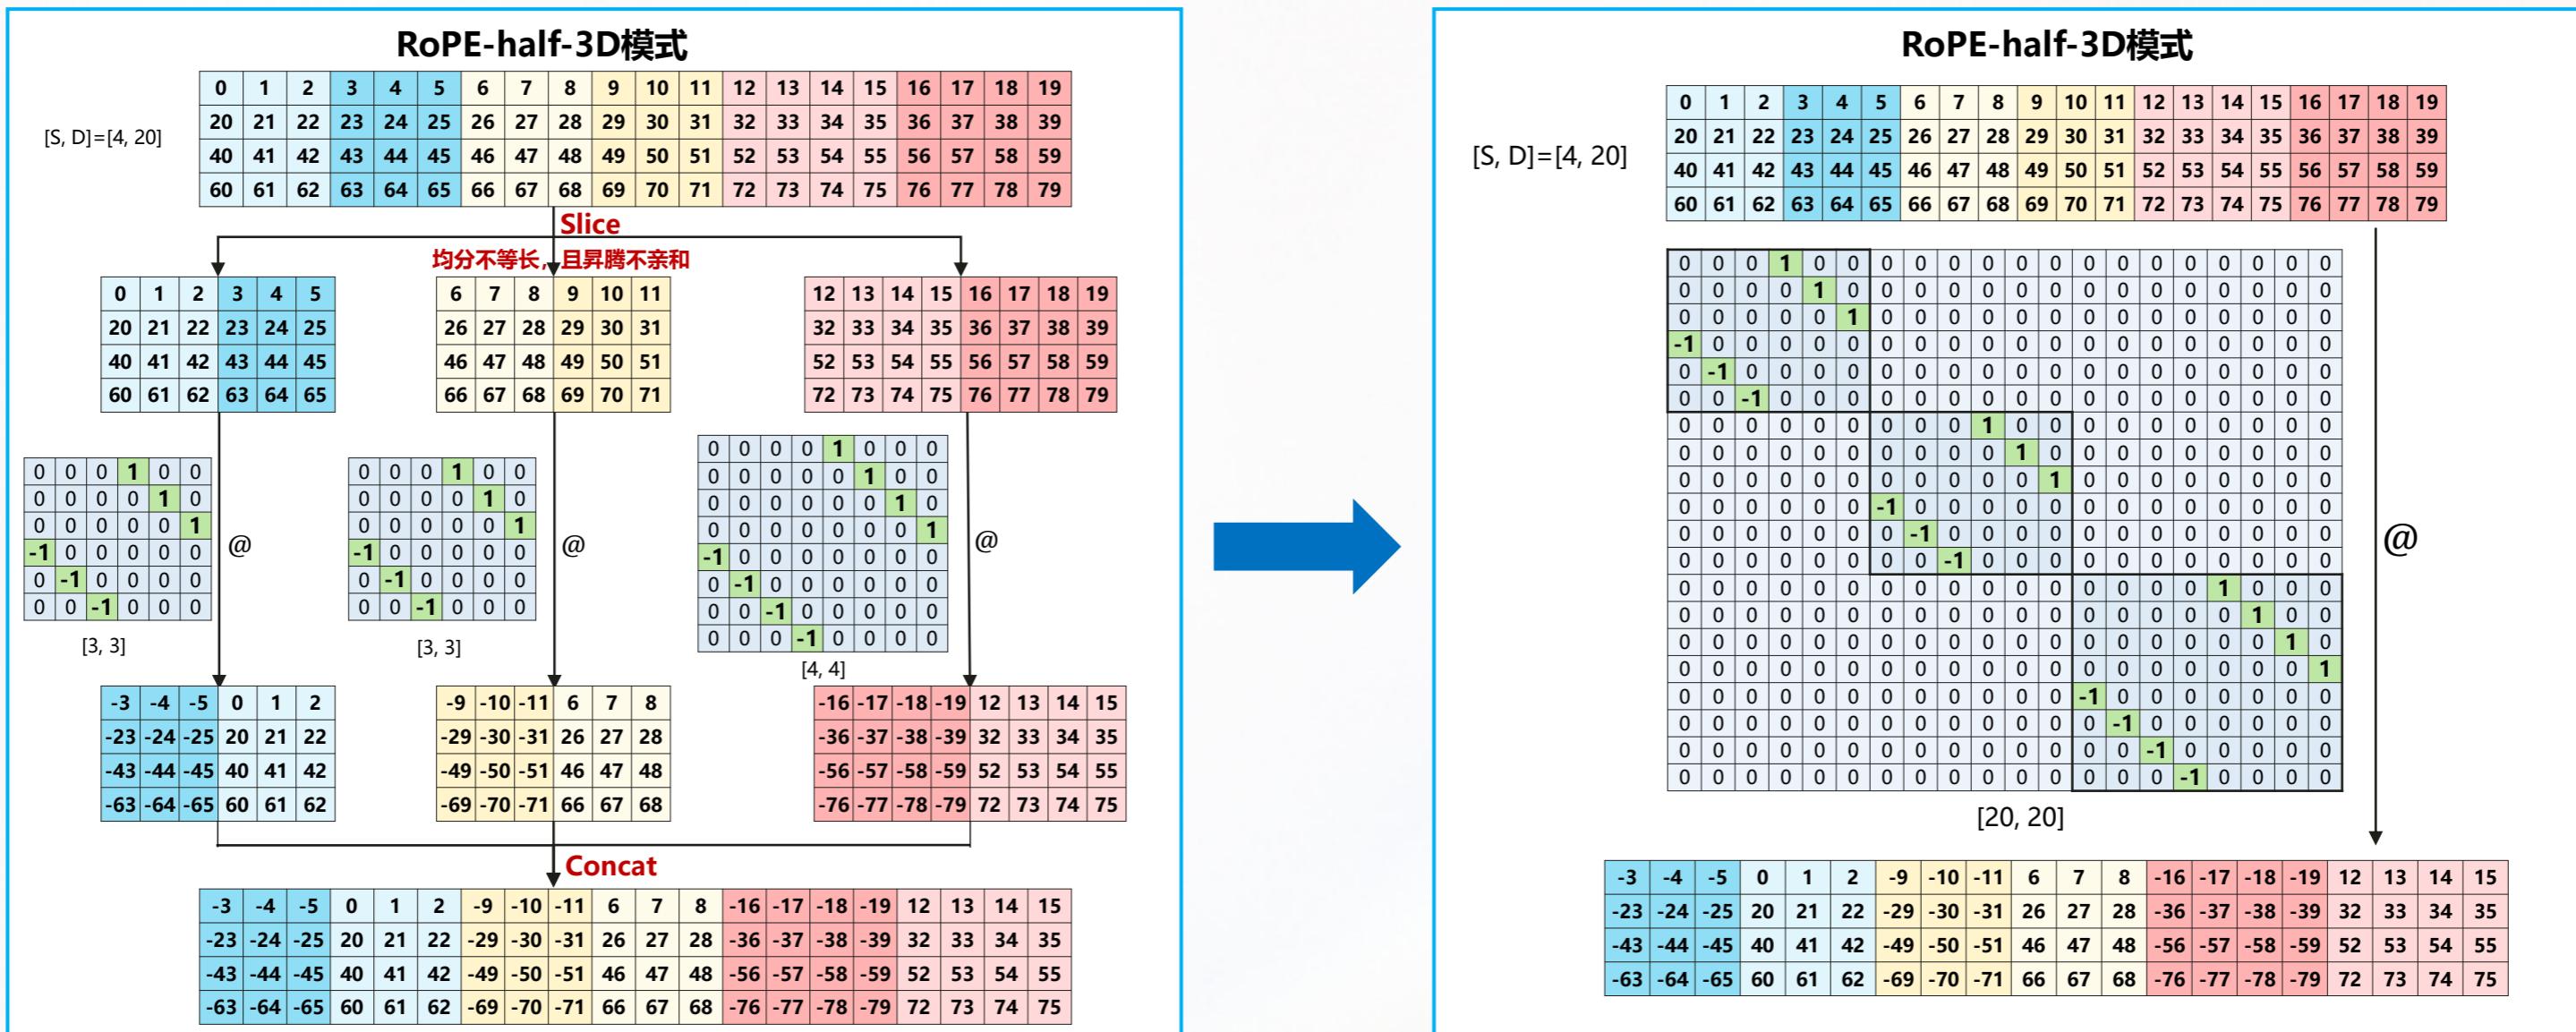
RoPE-Matrix: 昇腾亲和的RoPE算法实现

Step1: 以Cube计算替代Vector计算的新RoPE实现: 针对原RoPE中存在大量的对数据切分与合并操作，通过构造对应的操作矩阵以矩阵乘实现相关功能，以此替代昇腾不亲和的Vector操作。



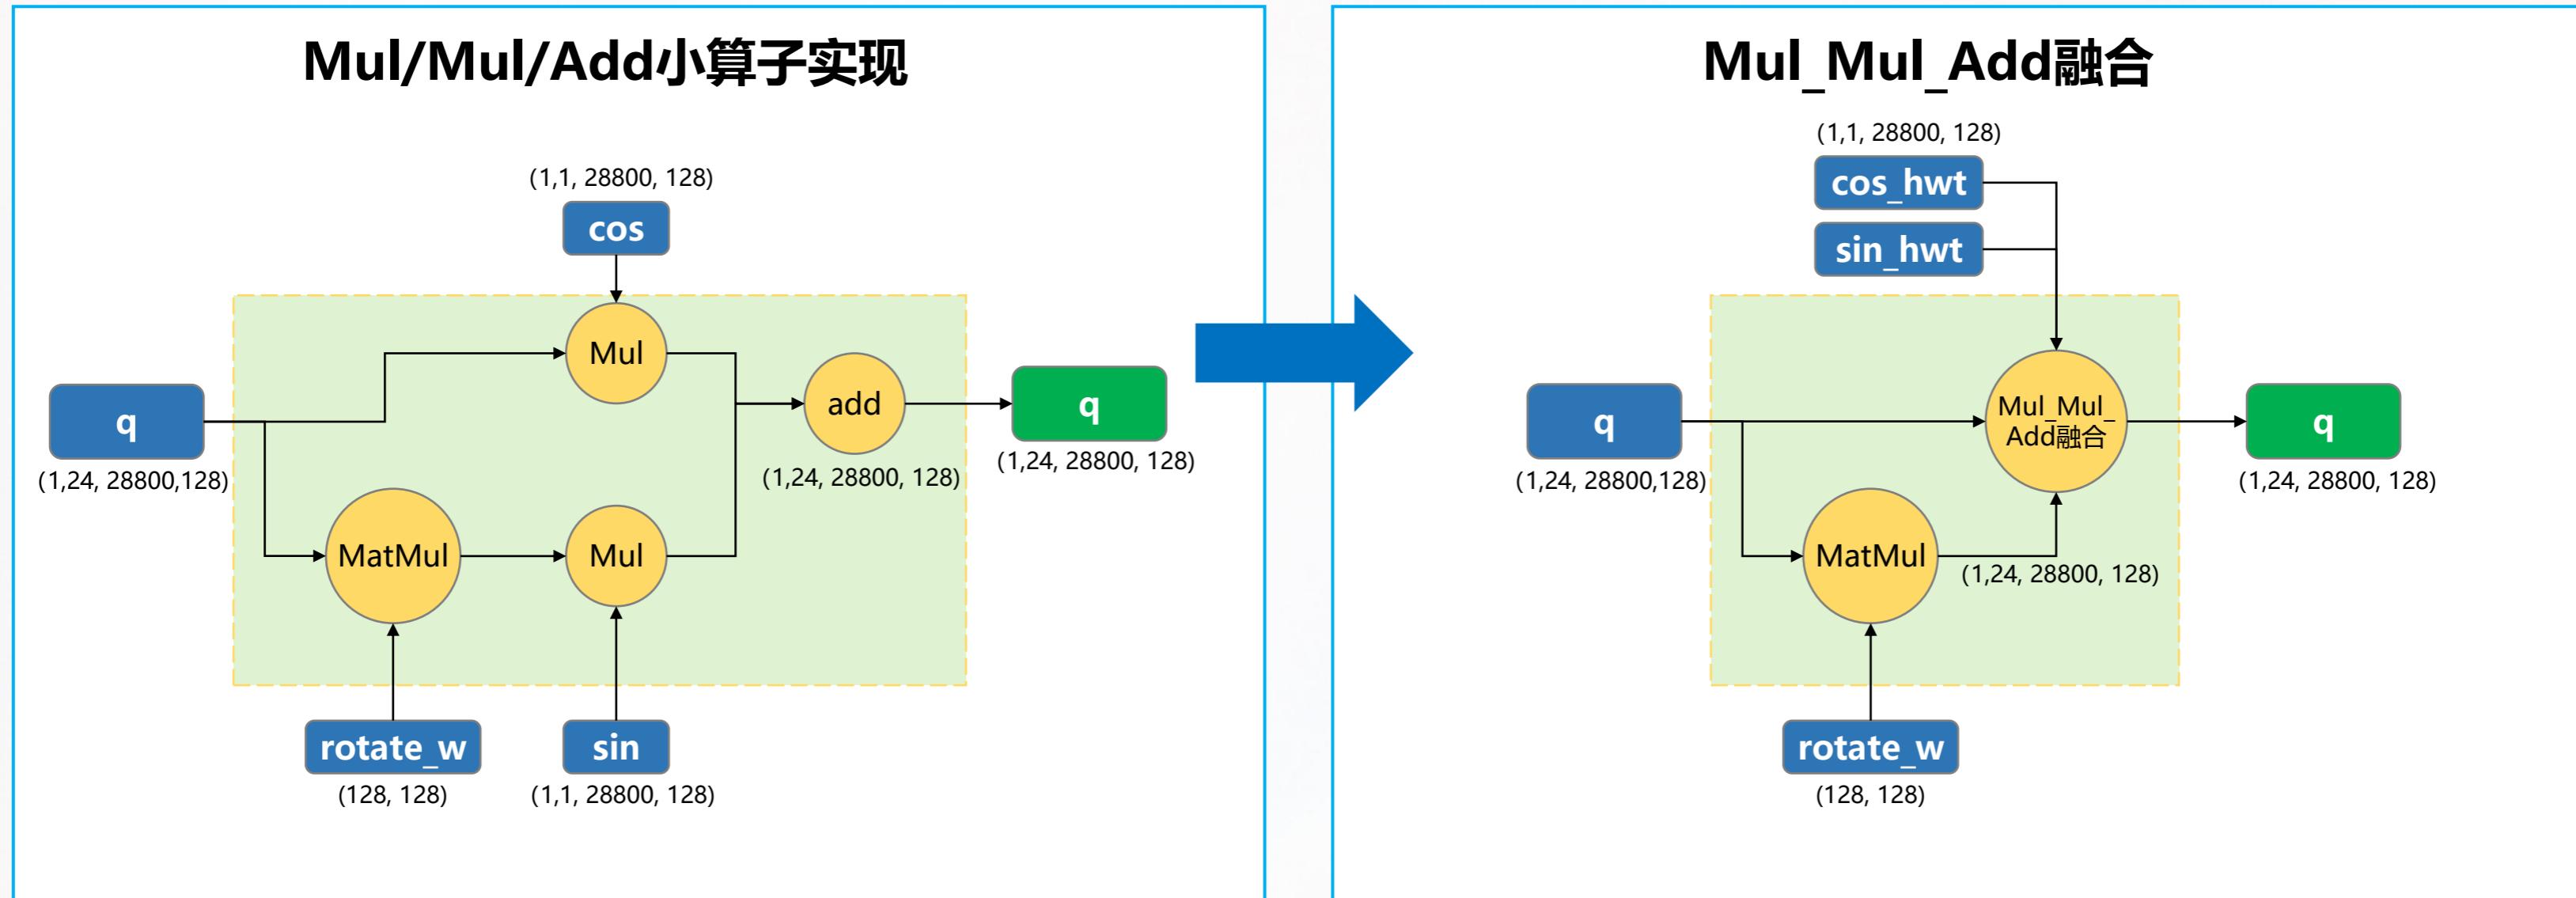
RoPE-Matrix: 昇腾亲和的RoPE算法实现

Step2: RoPE-2D/3D模式下，不同数据的Cube Rotate可以合并为更大的rotate_mat，一次性实现完整输入的全部计算操作。



RoPE-Matrix: 昇腾亲和的RoPE算法实现

Step3: 对Vector计算实现编译融合, 进一步实现加速。

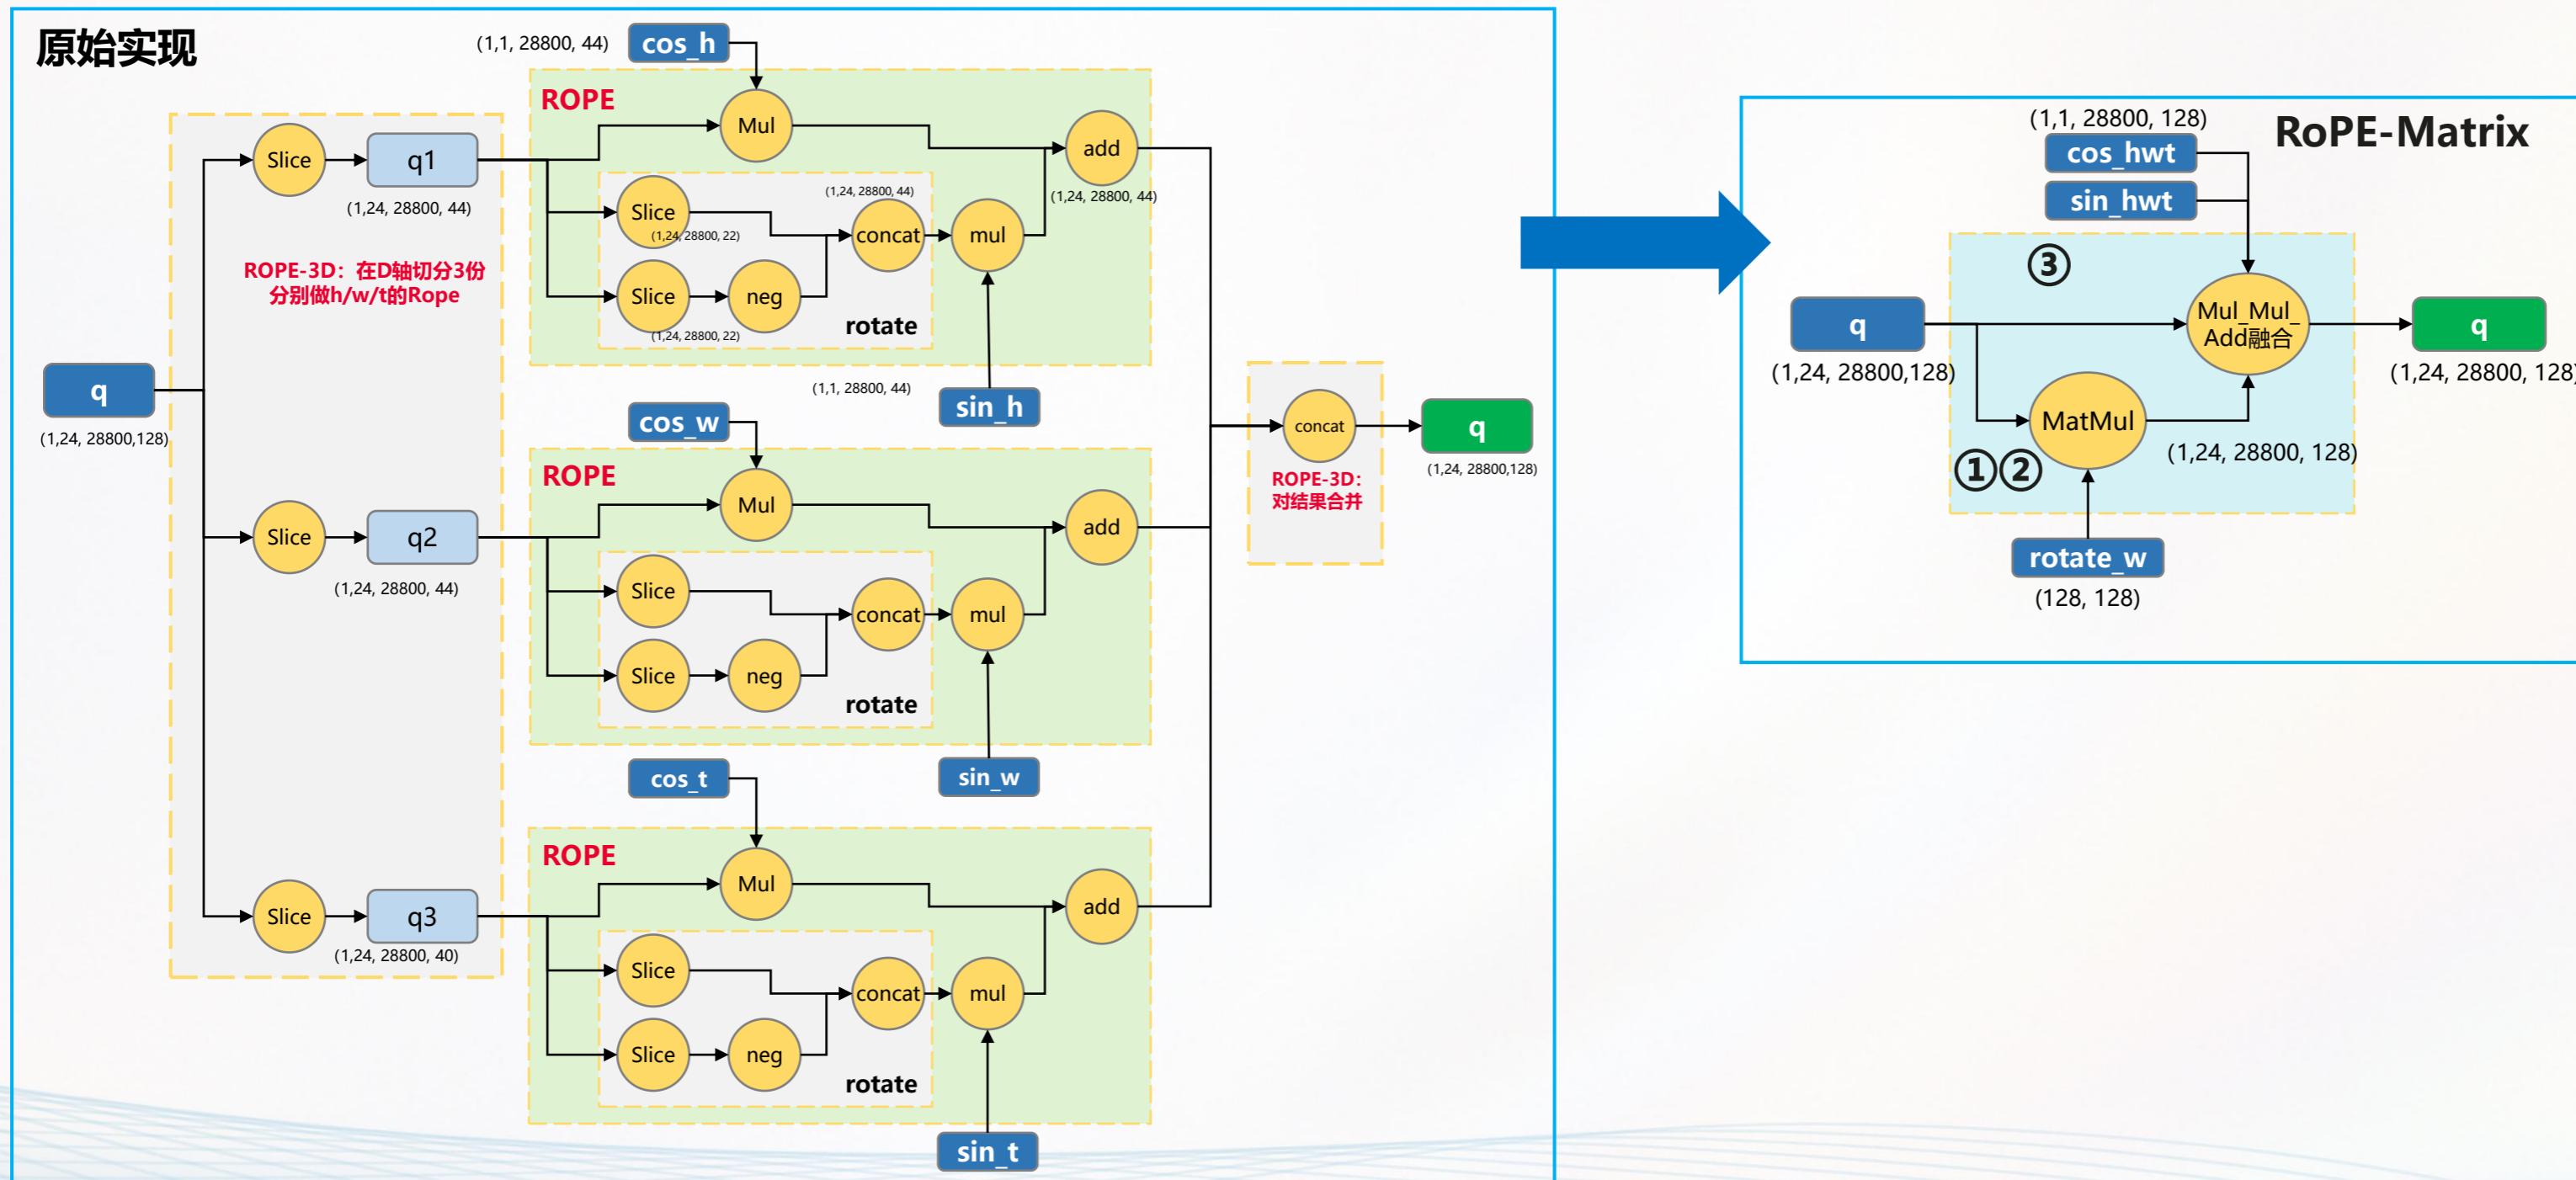


RoPE-Matrix: 昇腾亲和的RoPE算法实现

Step1: 以Cube计算替代Vector计算的新RoPE实现;

Step2: RoPE-2D/3D模式下, 用大旋转矩阵替换Slice-Concat操作;

Step3: 引入torch.compile对Vector计算实现编译融合。



RoPE-Matrix：昇腾亲和的RoPE算法实现

单算子收益：相比小算子实现均有显著性正收益（平均提速300%），interleave（平均提速130%）和half场景（267%）下相比融合大算子也有明显收益，覆盖绝大部分典型网络。

模式 bf16, [1, 24, 28800, 128]		典型网络	小算子实现 (无约束) 耗时: ms	CANN大算子 (约束条件) 耗时: ms	RoPE-Matrix 耗时: ms	相比 小算子	相比 CANN大算子
interleave	RoPE-1D dim=128	DeepSeek-V3/R1 LLaMA1/2/3/4 Qwen1/2/2.5	2.3	1.0 (其中2D/3D按合 并成整体高效执行)	0.8	2.9x	1.3x
	RoPE-2D dim=64+64	Qwen-Image Flux PG-T2I	2.6			3.3x	
	RoPE-3D (dim=44+44+40)	Wan2.1/2.2 hunyuanyvideo PG-T2V/I2V	2.8			3.6x	
half	RoPE-1D dim=128	PG-72BA3	2.1	0.35	0.8	2.6x	0.4x (Memory Bound+CV串行)
	RoPE-2D dim=64+64	-	2.5	1.1 含Slice+Concat		3.1x	1.4x
	RoPE-3D (dim=44+44+40)	StepVideo/OpenSORA	2.9	5.0 含Slice+Concat		3.6x	6.2x

RoPE-Matrix: 昇腾亲和的RoPE算法实现

整网收益: a) 开源网络推理, LLM/多模态理解/多模态生成, 分别实现4.7%/4.5%/4.5%加速

Task	Model	Parameters	QKShape[BNSD]	Baseline(ms)	RoME(ms)	Speedup
LLM	Llama3.1	8B	[1, 24, 8192, 128]	50.8	48.1	4.7%
VLM	internvl3	8B	[1, 28, 1846, 128]	174	166	4.5%
Edit	FLUX.1-Kontext	12B	[1, 24, 8704, 128]	1059	978	8.3%
T2I	Qwen-Image	20B	[1, 24, 6032, 128]	1457	1401	4%
T2V	HunyuanVideo	13B	[1, 24, 42840, 128]/540P	14410	14400	0.1%
			[1, 24, 8192, 128]/720P	40120	40100	0.1%
T2V	Wan2.2	1.3B	[1, 12, 6630, 128]/272P	209	199	3.9%
			[1, 12, 17550, 128]/480P	790	759	3.4%
		14B	[1, 40, 6630, 128]/272P	1170	1120	4.3%
			[1, 40, 17550, 128]/480P	4500	4350	3.3%

b) 开源网络训练, 多模态生成 (wan2.2-1.3B), 分别实现4%/2.3%加速

Inputs	RoPE (s)	RoME (s)	Speedup
272P	1.27	1.22	4%
480P	3.37	3.30	2.3%

RoPE-Matrix：昇腾亲和的RoPE算法实现

强泛化性：新方案支持任意多维RoPE的不同mode，且相比融合大算子，没有任何约束条件

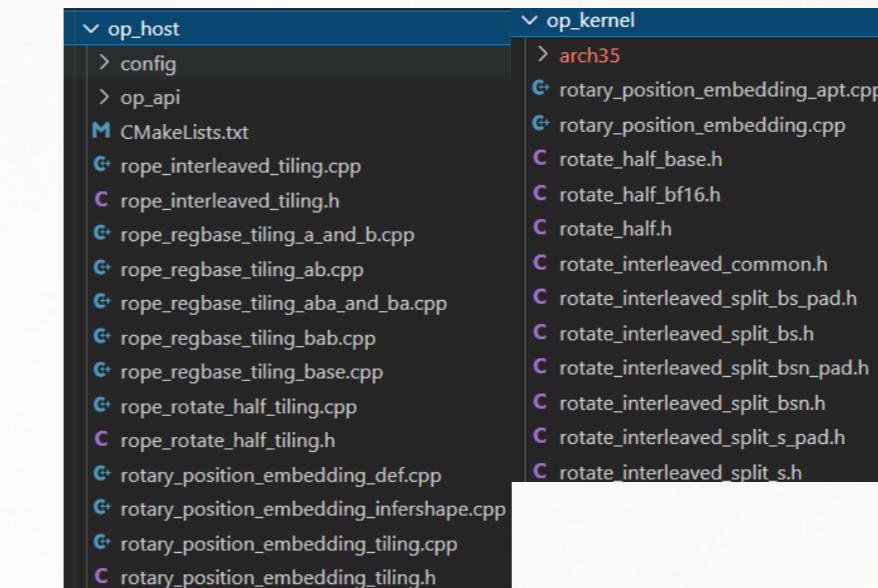
torch_npu.npu_rotary_mul约束条件

约束说明

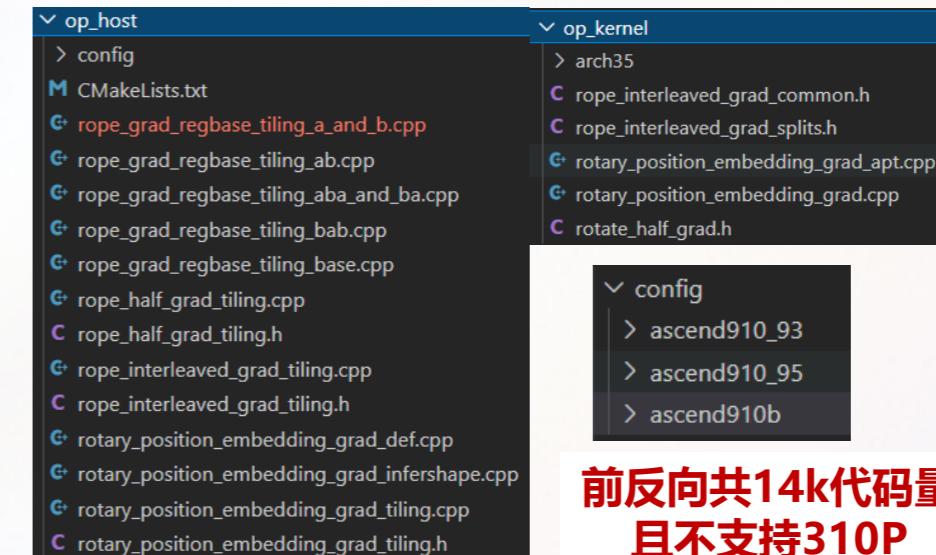
- **jit_compile=False**场景（适用Atlas A2 训练系列产品^[*], Atlas A3 训练系列产品^[*]）：
 - half模式：
 - input: layout支持: *BNSD*、*BSND*、*SBND*; $D < 896$, 且为2的倍数; $B, N < 1000$; 当需要计算*cos/sin*的反向梯度时, $B * N <= 1024$ 。
 - r1、r2: 数据范围: [-1, 1]; 对应input layout的支持情况:
 - x为*BNSD*: *11SD*、*B1SD*、*BNSD*;
 - x为*BSND*: *1S1D*、*BS1D*、*BSND*;
 - x为*SBND*: *S11D*、*SB1D*、*SBND*。
 - 预知:
*half*模式下, 当输入*layout*是*BNSD*, 且*D*为非32Bytes对齐时, 建议不使用该融合算子（模型启动脚本中不开启--use-fused-roty-pos-emb选项）, 否则可能出现性能下降。
- interleave模式:
 - input: layout支持: *BNSD*、*BSND*、*SBND*; $B * N < 1000$; $D < 896$, 且*D*为2的倍数;
 - r1、r2: 数据范围: [-1, 1]; 对应input layout的支持情况:
 - x为*BNSD* : *11SD*;
 - x为*BSND* : *1S1D*;
 - x为*SBND* : *S11D*

- **jit_compile=True**场景（适用Atlas 训练系列产品^[*], Atlas A2 训练系列产品^[*], Atlas 推理系列产品^[*]）：
- 仅支持*rotary_mode*为*half*模式, 且r1/r2 layout一般为*11SD*、*1S1D*、*S11D*。
- shape要求输入为4维, 其中*B*维度和*N*维度数值需小于等于1000, *D*维度数值为128。

CANN RotaryPositionEmbedding前向实现



CANN RotaryPositionEmbedding反向实现



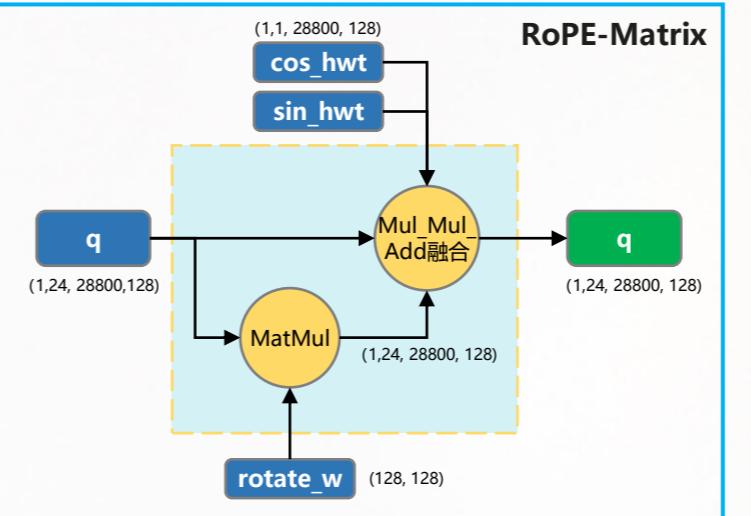
前反向共14k代码量
且不支持310P

RoPE-Matrix: 昇腾亲和的RoPE算法实现

- **RoPE-Matrix**

- 一种昇腾亲和的以Cube计算替代部分Vector计算的**RoPE矩阵乘**实现方式；
- 算子维度取得CANN小算子**300%**加速、CANN大算子**148%**加速；
- 整网收益：a) 开源网络推理，LLM/多模态理解/多模态生成，分别实现**4.7%/4.5%/4.5%**加速；b) 开源网络训练，多模态生成（wan2.2-1.3B），分别实现4%/2.3%加速
- 高泛化性，同时覆盖interleave/half/quarter/interleave-half等模式
- 论文已投CVPR 2026

$$output = x * \cos + (x @ rotate_mat) * \sin$$



This screenshot shows the OpenReview.net interface for a paper titled "Efficient Matrix Implementation for Rotary Position Embedding". The paper was submitted on 06 Nov 2025 and modified on 14 Nov 2025. It is a Conference Submission under the Senior Area Chairs, Area Chairs, Reviewers, Authors category. The subject area is "Deep learning architectures and techniques". Key words include "rotation embedding, rotary position embedding, efficient". The abstract discusses the inefficiencies of existing RoPE implementations and how RoME (Rotary Position Embedding Matrix) improves them by using matrix operations instead of vector operations. The paper has been withdrawn.

此处预留直播画面LOGO
不放置内容

RoPE-Matrix: 昇腾亲和的RoPE算法实现

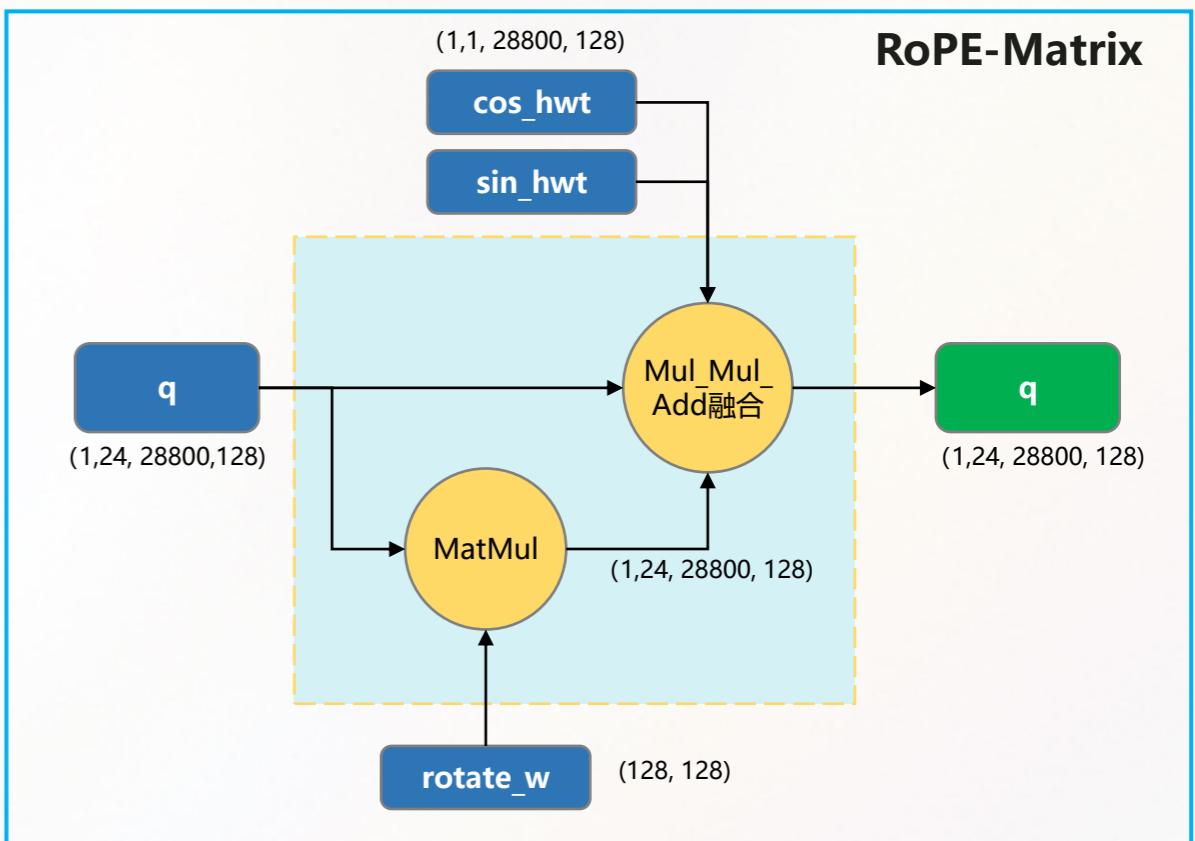
- **RoPE-Matrix**

- 一种昇腾亲和的以Cube计算替代部分Vector计算的**RoPE矩阵乘**实现方式；
- 算子维度取得CANN小算子**300%**加速、CANN大算子**148%**加速；
- 整网收益：a) 开源网络推理，LLM/多模态理解/多模态生成，分别实现**4.7%/4.5%/4.5%**加速；b) 开源网络训练，多模态生成(wan2.2-1.3B)，分别实现4%/2.3%加速
- 高泛化性，同时覆盖interleave/half/quarter/interleave-half等模式
- 论文已投CVPR 2026

The screenshot shows a paper submission page on OpenReview.net. The title is "Efficient Matrix Implementation for Rotary Position Embedding". Key details include:

- Submitted on 06 Nov 2025 (modified: 14 Nov 2025)
- CVPR 2026 Conference Submission
- Conference, Senior Area Chairs, Area Chairs, Reviewers, Authors
- CC BY 4.0
- Subject Area: Deep learning architectures and techniques
- Keywords: position embedding, rotary position embedding, efficient
- Student Paper: No
- External Links: A confirmation message about no external links.
- Abstract: A detailed description of RoPE-Matrix, mentioning its efficiency and how it replaces vector operations with unified matrix transformations.
- Paper Submission: PDF file
- Compute Report: PDF file
- Submission Number: 18746

$$output = x * \cos + (x @ rotate_mat) * \sin$$



RoPE-Matrix：融合算子 开发与贡献

02

RoPE-Matrix：融合算子开发与贡献

代码开发：基于Ascend C示例代码及CANN RotaryPositionEmbedding开源仓代码，共3人周完成RoPE-Matrix算子开发，精度性能达标。



开源代码参考

算子编译工程参考：

https://gitee.com/ascend/samples/tree/master/operator/ascendc/0_introduction/13_matmulleakyrelu_kernellaunch

MIX算子框架及MatMul实现参考：

https://gitee.com/ascend/samples/blob/master/operator/ascendc/0_introduction/22_baremix_kernellaunch/

Vector实现参考CANN内置算子

https://gitcode.com/cann/ops-transformer/tree/master/poseembedding/rotary_position_embedding

```
KERNEL_TASK_TYPE_DEFAULT(KERNEL_TYPE_MIX_AIC_1_2);
if ASCEND_IS_AIC {
    ...
    AscendC::CrossCoreSetFlag<0x2, PIPE_FIX>(0x8); // set flag after matmul
}
if ASCEND_IS_AIV {
    AscendC::CrossCoreWaitFlag(0x8); // wait matmul outputs
...
}
```

PYBIND算子调用方式

```
at::Tensor rope_matrix_kernel_bf16(at::Tensor x, at::Tensor y, at::Tensor sin, at::Tensor cos)
{
    ...
    ACLRT_LAUNCH_KERNEL(rope_matrix_kernel_bf16)(
        blockDim, aclstream,
        (uint8_t*)(x.storage().data()),
        (uint8_t*)(y.storage().data()),
        (uint8_t*)(sin.storage().data()),
        (uint8_t*)(cos.storage().data()),
        (uint8_t*)(output.storage().data()),
        (uint8_t*)(workspace_tensor.storage().data()),
        tilingDevice);
}
```

此处预留直播画面LOGO
不放置内容

RoPE-Matrix：融合算子开发与贡献

代码开发：基于Ascend C示例代码及CANN RotaryPositionEmbedding开源仓代码，共3人周完成RoPE-Matrix算子开发，精度性能达标。



开源代码参考

算子编译工程参考：

https://gitee.com/ascend/samples/tree/master/operator/ascendc/0_introduction/13_matmulleakyrelu_kernellaunch

MIX算子框架及MatMul实现参考：

https://gitee.com/ascend/samples/blob/master/operator/ascendc/0_introduction/22_ba_remix_kernellaunch/

Vector实现参考CANN内置算子

https://gitcode.com/cann/ops-transformer/tree/master/poseembedding/rotary_position_embedding

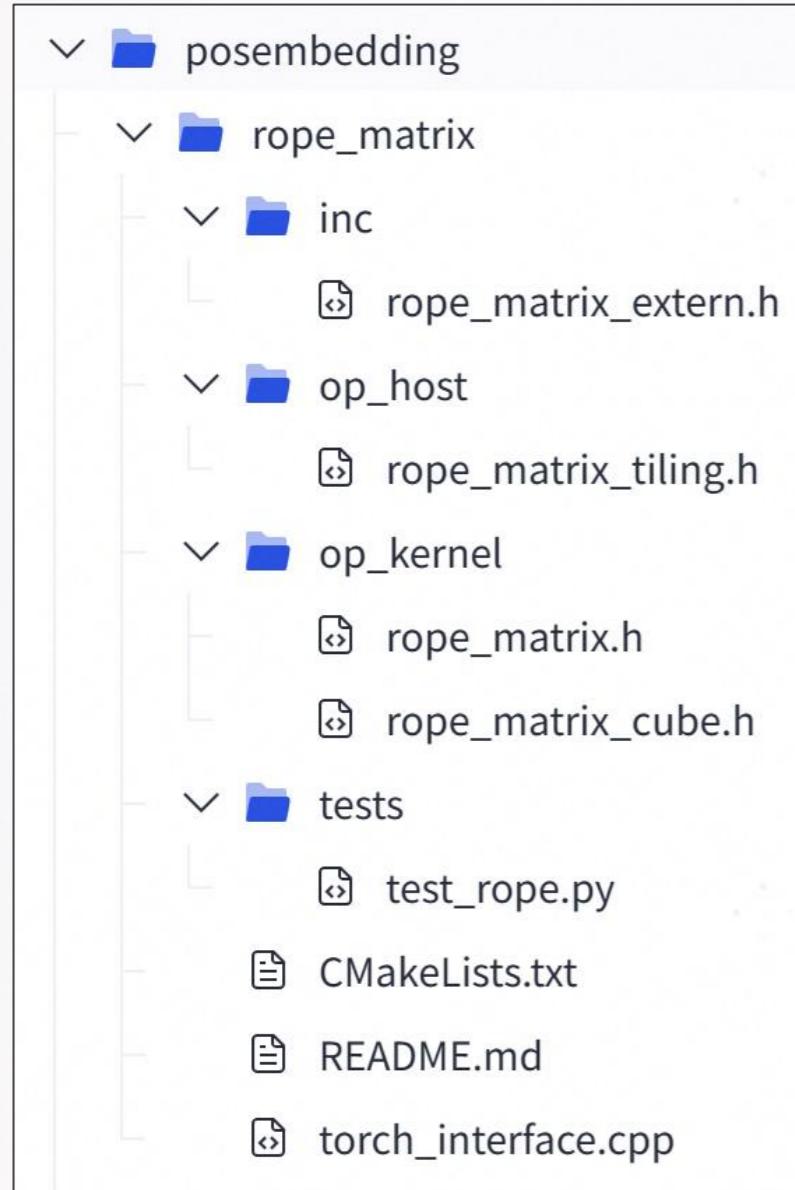
KERNEL_TASK_TYPE_DEFAULT(KERNEL_TYPE_MIX_AIC_1_2);
if ASCEND_IS_AIC {
 ...
 AscendC::CrossCoreSetFlag<0x2, PIPE_FIX>(0x8); // set flag after matmul
}
if ASCEND_IS_AIV {
 AscendC::CrossCoreWaitFlag(0x8); // wait matmul outputs
}
...

PYBIND算子调用方式

```
at::Tensor rope_matrix_kernel_bf16(at::Tensor x, at::Tensor y, at::Tensor sin, at::Tensor cos)  
{  
    ...  
  
    ACLRT_LAUNCH_KERNEL(rope_matrix_kernel_bf16)(  
        blockDims, aclstream,  
        (uint8_t*)(x.storage().data()),  
        (uint8_t*)(y.storage().data()),  
        (uint8_t*)(sin.storage().data()),  
        (uint8_t*)(cos.storage().data()),  
        (uint8_t*)(output.storage().data()),  
        (uint8_t*)(workspace_tensor.storage().data()),  
        tilingDevice);  
}
```

RoPE-Matrix: 融合算子开发与贡献

算子开发: 已开源至https://gitcode.com/cann/ops-transformer/tree/master/experimental/poseembedding/rope_matrix



Tiling关键代码

```
optiling::TCubeTiling tilingData;
auto ascendcPlatform = platform_ascendc::PlatformAscendCManager::GetInstance(socVersion);
matmul_tiling::MultiCoreMatmulTiling tilingApi(*ascendcPlatform);

tilingApi.SetDim(usedCoreNum);
tilingApi.SetAType(leftPosition, leftFormat, leftDtype, isTransA);
tilingApi.SetBType(rightPosition, rightFormat, rightDtype, isTransB);
tilingApi.SetCType(resultPosition, resultFormat, resultDtype);

tilingApi.SetOrgShape(M, N, K); // 完成的MNK大小, 单位为元素个数
tilingApi.SetShape(M, N, K); // matmul计算形状的MNK, 考虑脏数据
tilingApi.SetSingleShape(calSingleCoreM, baseSize, baseSize);
tilingApi.SetFixSplit(baseM, baseN, -1);
tilingApi.SetBias(isBias);
tilingApi.SetBufferSpace(-1, -1, -1);

int64_t res = tilingApi.GetTiling(tilingData);
```

MatMul计算关键代码

```
int total_N = B * H;
for (int Nindex = 0; Nindex < total_N; ++Nindex) {
    int offsetA = 0, offsetB = 0, offsetC = 0;
    int tailM = 0, tailN = 0;
    bool isTransA = false, isTransB = false;

    CalcGMOFFset(AscendC::GetBlockIdx(), offsetA, offsetB, offsetC,
                  tailM, tailN, B, H, M, N, K, Nindex);

    matmulObj.SetTensorA(aGlobal[offsetA], isTransA);
    matmulObj.SetTensorB(bGlobal[offsetB], isTransB);
    matmulObj.SetTail(tailM, tailN);
    matmulObj.IterateAll(cGlobal[offsetC]);
}

matmulObj.End();
```

Vector计算关键代码

```
_aicore_ inline void SingleStepProcess(uint32_t progress, uint64_t copyLength, uint64_t calcLength) {
    uint64_t xOffset;
    uint64_t rOffset;
    uint64_t bnLoopXStartOffset;
    uint64_t progressOffset;
    uint64_t batchOffset;
    rOffset = progress * this->sdBlockSize;
    CopyInR(rOffset, copyLength);
    LocalTensor<T> cosLocal = inQueueCos.DeQue<T>();
    LocalTensor<T> sinLocal = inQueueSin.DeQue<T>();
    LocalTensor<float> cosFp32 = cosBuf.Get<float>();
    LocalTensor<float> sinFp32 = sinBuf.Get<float>();

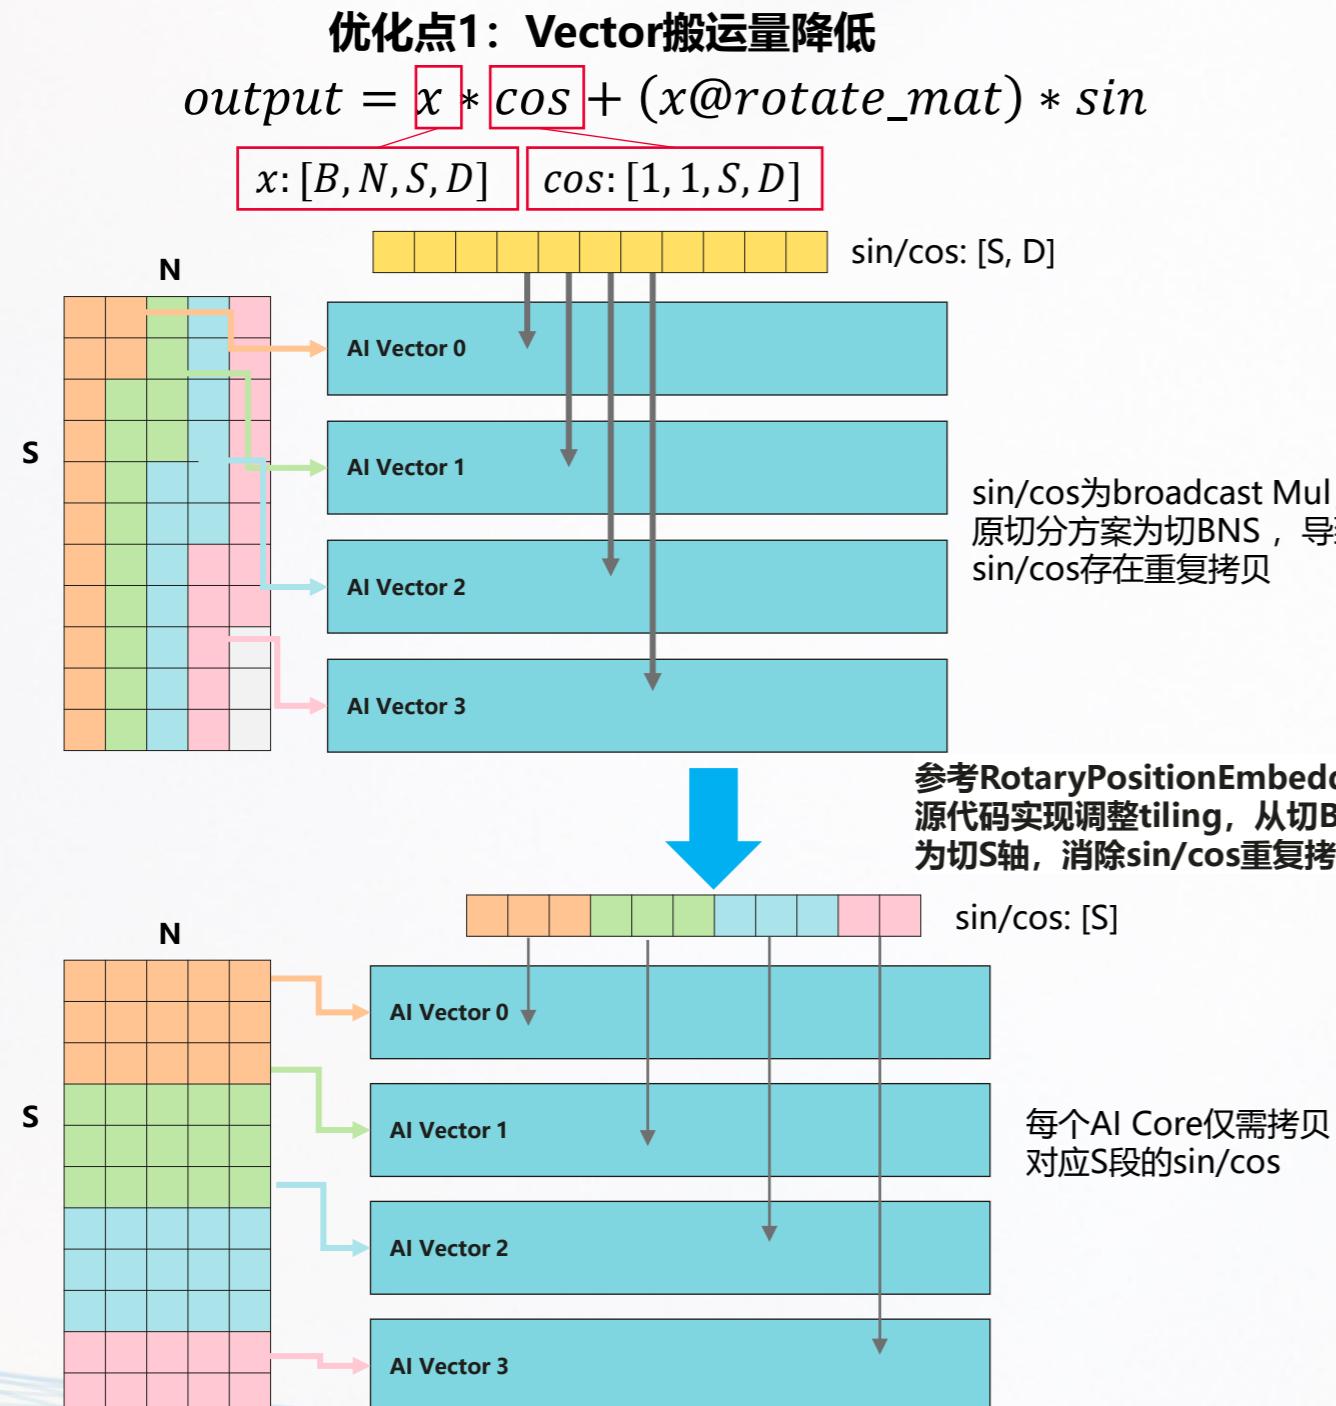
    Cast(cosFp32, cosLocal, RoundMode::CAST_NONE, calcLength);
    Cast(sinFp32, sinLocal, RoundMode::CAST_NONE, calcLength);
    inQueueCos.FreeTensor<T>(cosLocal);
    inQueueSin.FreeTensor<T>(sinLocal);

    bnLoopXStartOffset = progress * this->sdBlockSize;
    for (uint32_t bnLoop = 0; bnLoop < this->bnSize; bnLoop++) {
        xOffset = bnLoopXStartOffset + bnLoop * this->sdSizeTotal;
        CopyInX(xOffset, copyLength);
        Compute(cosFp32, sinFp32, calcLength);
        CopyOut(xOffset, copyLength);
    }
}

_aicore_ inline void ComputeInner(LocalTensor<float>& x, LocalTensor<float>& xNew,
                                  LocalTensor<float>& cos, LocalTensor<float>& sin,
                                  uint32_t calcLength) {
    Mul(x, x, cos, calcLength);
    Mul(xNew, xNew, sin, calcLength);
    Add(xNew, xNew, x, calcLength);
}
```

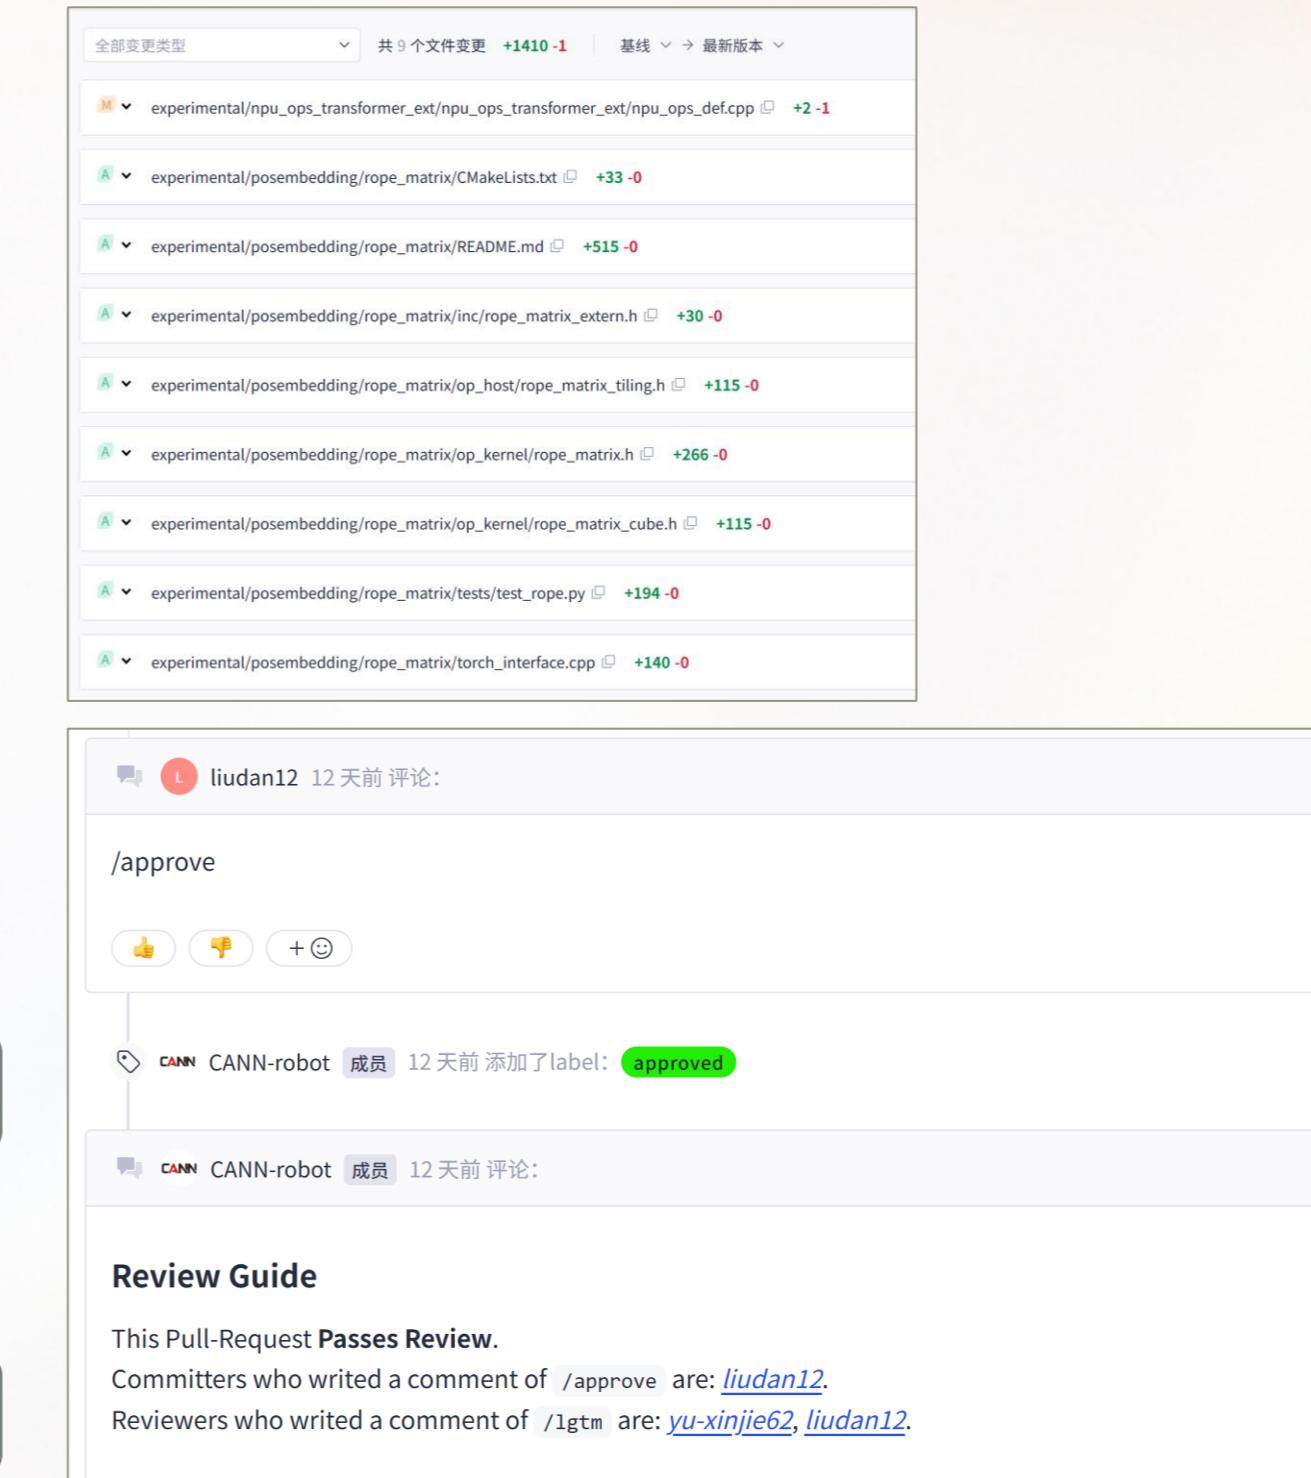
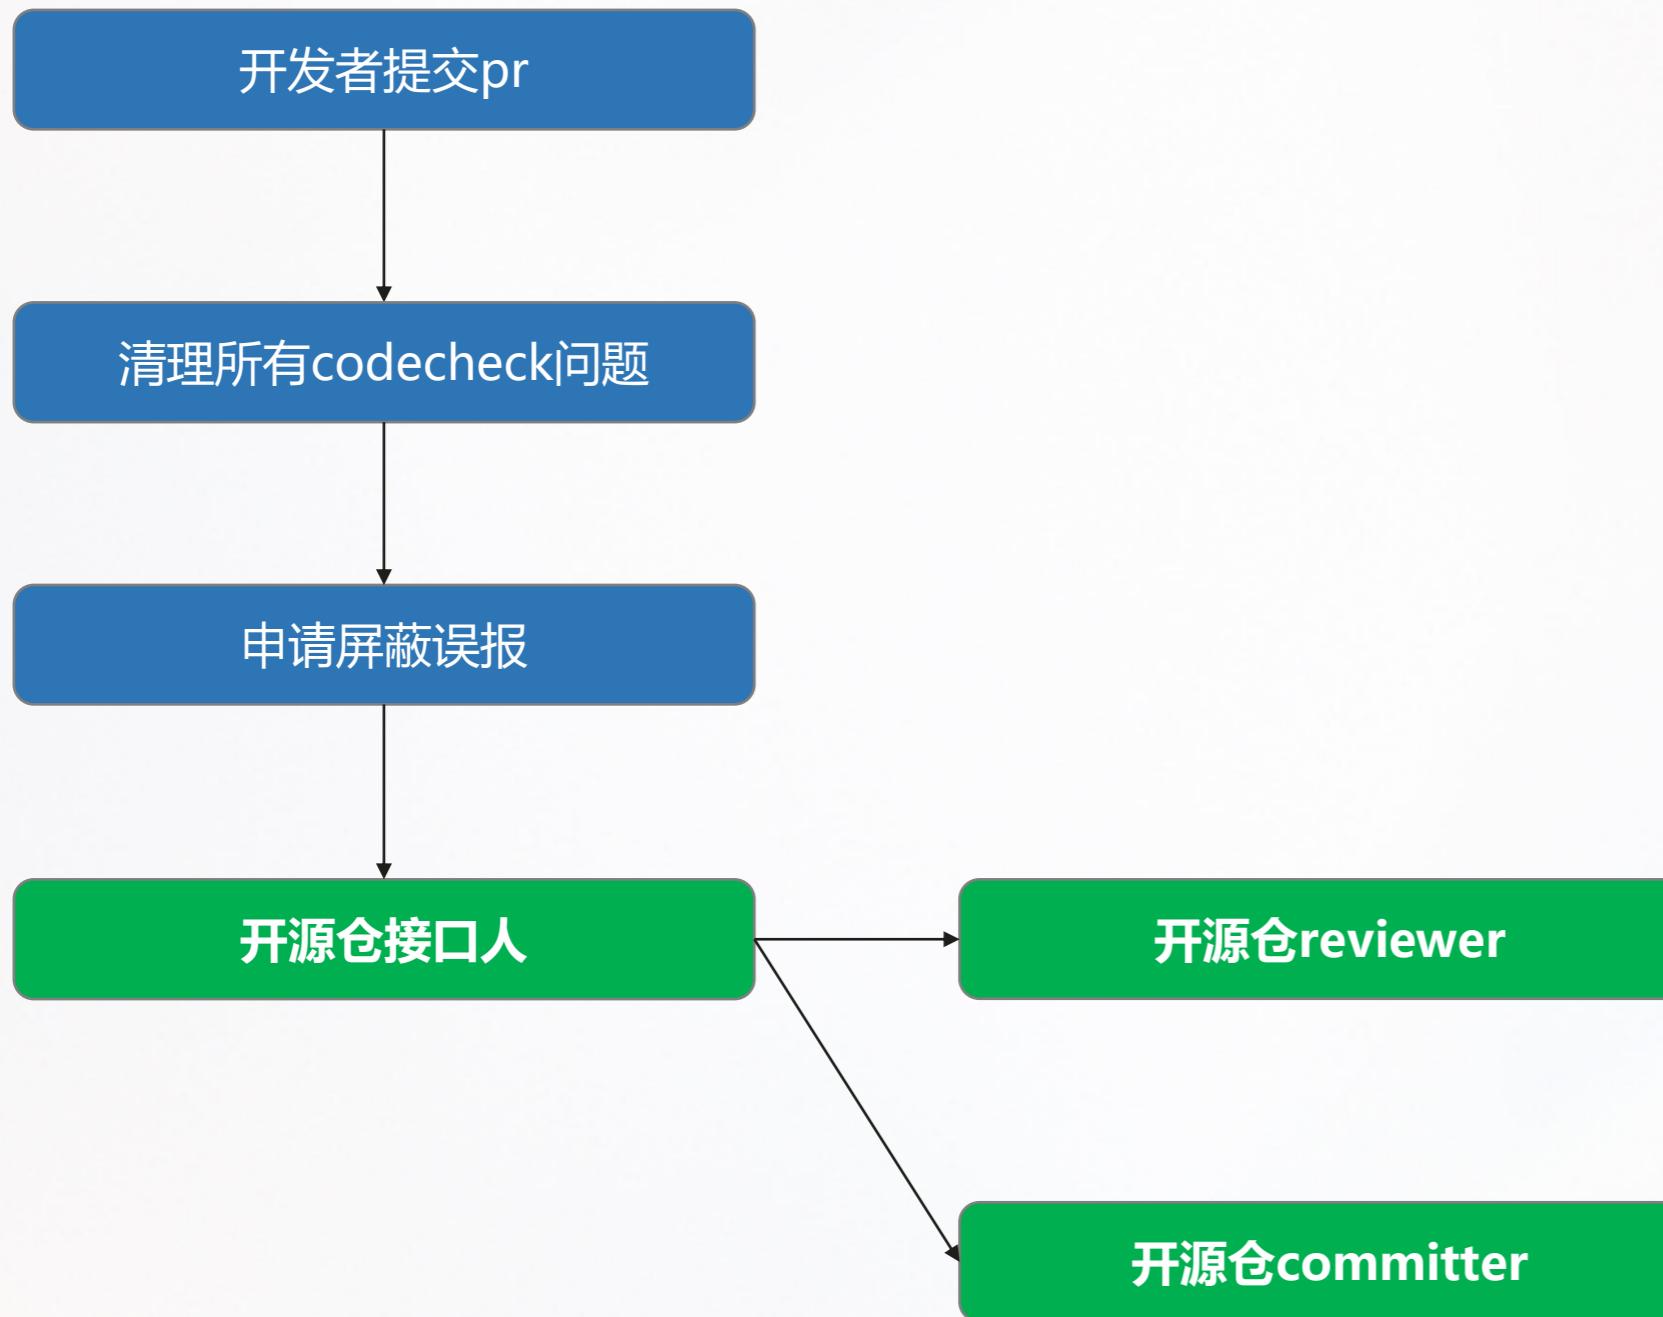
RoPE-Matrix：融合算子开发与贡献

性能优化，单算子耗时从10ms优化至0.8ms达到预期目标，达成memory bound，优化方法：a) 参考RotaryPositionEmbedding开源代码实现调整tiling，降低Vector搬运量；b) MatMul右矩阵大小仅32KB，设置常驻L1，降低Cube搬运量



RoPE-Matrix: 融合算子开发与贡献

代码合入: 核心代码量0.7k, 整体合入过程顺利, 约1天。



RoPE-Matrix：融合算子开发与贡献

算子源码：https://gitcode.com/cann/ops-transformer/tree/master/experimental/poseembedding/rope_matrix

算子编译与安装：https://gitcode.com/cann/ops-transformer/tree/master/experimental/npu_ops_transformer_ext

1. 进入目录，安装依赖

```
cd experimental/npu_ops_transformer_ext  
pip install -r requirements.txt
```

2. 从源码构建.whl包

```
python -m build --wheel -n
```

3. 安装构建好的.whl包

```
pip install dist/xxx.whl
```

算子调用：https://gitcode.com/cann/ops-transformer/blob/master/experimental/poseembedding/rope_matrix/tests/test_rope.py

```
import torch  
import torch_npu  
import npu_ops_transformer_ext
```

```
B, N, S, D = 1, 24, 28800, 128  
dtype = torch.bfloat16  
x = torch.randn((B, N, S, D), dtype=dtype).npu()  
mat = get_interleave_matrix(dim=D).to(dtype).npu()  
sin = torch.randn((1, 1, S, D), dtype=dtype).npu()  
cos = torch.randn((1, 1, S, D), dtype=dtype).npu()  
out = torch.ops.npu_ops_transformer_ext.rope_matrix(x, mat, sin, cos)
```

RoPE-Matrix：融合算子开发与贡献

RoPE-Matrix首个合入ops-transformer项目库的算子

A screenshot of a GitHub pull request page. The title is "add new operator ropematrix #299". It shows the pull request has been merged by "xuyun15" and is now part of the "cann/ops-transformer: master" branch. The pull request has 41 discussions, 1 submission, 0 reviews, and 9 file changes. The author is listed as "xuyun15". The "Description" section notes that the new rope matrix algorithm is mathematically equivalent to the existing rope implementation. The "Related Issues" section lists issue #299 as a requirement for adding an interleaved rope matrix operator.

算子开发流程及经验总结文档发布于社区

A screenshot of an Ascend developer blog post. The title is "[开源贡献案例] 端到端打通transformer仓experimental路径首个开源mix算子". The post discusses the development process and experience of contributing to the experimental path of the transformer repository. It includes sections on background, requirements, compilation errors, and namespace conflicts. The post is categorized under "新人帖" and "算子". It was published by "xuyun15" on November 24, 2023, with 452 likes and 13 comments.

共提出10个issue，其中9个已被accept，6个已关闭

Issue编号	内容	是否accept	当前状态
61	[Requirement 需求建议]: 新增interleave模式优化rope算子	Y	close
119	[Documentation 文档反馈]: Matmul高阶API性能调优文档缺失	Y	close
162	[Requirement 需求建议]: 算子贡献默认规则未描述清晰	Y	close
184	[Documentation 文档反馈]: experimental路径下使用RunOpApiV2完成算子适配调用缺乏详细信息	Y	close
187	[Documentation 文档反馈]: experimental路径下编译依赖说明不全	Y	close
195	[Documentation 文档反馈]: 官网matmul高阶api的约束没有说明清楚	N	open
198	[Documentation 文档反馈]: experimental路径下一些编程规范没有说清楚	Y	open
203	[Documentation 文档反馈]: 特定场景下matmul引发精度问题，文档未提及	Y	open
206	[Documentation 文档反馈]: 高阶api对ASCENDC_CUBE_ONLY描述的不够清晰，与用例使用不一致	Y	close
230	[Requirement 需求建议]: experimental路径下部分编译能力提供	Y	open

ops-transformers仓库discussion经验沉淀

A screenshot of a GitHub discussion thread. The title is "使用<<>>方式实现算子FAQ #3". The post was created by "fanzijian" 8 days ago. It provides a reference document for implementing operators using the new API. The document includes sections on environment requirements, compilation errors, and namespace issues. It also links to a README file for more details.

Thanks!



访问CANN开源社区



关注昇腾CANN公众号

