

# CANN算子库ATVOSS

极简、高效、高性能、高扩展的编程范式

作者：李东锋/沈敏

时间：2025-12-03

<https://gitcode.com/cann>

**CANN**

# 目录

Part 1 ATVOSS

Part 2 ATVOSS开源仓

# 背景

## 当前现状

- Tiling计算繁杂
- 指令灵活,操作维度多
- 性能调优困难
- 感知芯片差异

## 我们的思考


- 如何帮助用户高效写出一个高性能算子

## 我们要做什么样的东西

- 它是极简的
- 它是高效的
- 它性能要有保障
- 它应该支持扩展

# 设计目标：极简、高效、高性能、高扩展,全面提升算子开发易用性

ATVOSS :Ascend C Templates for Vector Operator Subroutines



# 极简：采用极简表达式，编程界面简单易懂

Rms-norm计算公式：

$$\text{RmsNorm}(x_i) = \frac{x_i}{\text{Rms}(\mathbf{x})} g_i, \quad \text{where } \text{Rms}(\mathbf{x}) = \sqrt{\frac{1}{n} \sum_{i=1}^n x_i^2}$$

```
template<typename T>
struct RmsnormCompute {

    template <template <typename> class Tensor>
    __host__ __aicore__ constexpr auto Compute() const
    {
        auto x = PlaceHodler<1, Tensor<T>, Atvoss::Arguments::IN>();
        auto g = PlaceHodler<2, Tensor<T>, Atvoss::Arguments::IN>();
        auto y = PlaceHodler<3, Tensor<T>, Atvoss::Arguments::OUT>();
        auto tmp = PloceHolderTmpLike<1>(x);
        auto z = PloceHolderTmpLike<2, Pattern::AR>(x);
        return (tmp = x * x,
                z = ReduceSum<Pattern::AR>(tmp),
                z = Divs(z, GetShape<Int<-1>>(x)),
                z = Sqrt(z),
                tmp = Broadcast<Pattern::AB>(z),
                tmp = x / tmp,
                y = g * tmp);
    }
};
```

抽象化计算表达界面：

延迟计算

- Device/Kenrel/Block层只记录表达式，形成类型化抽象语法树
- Tile层对表达式进行实际计算操作

# 极简：精简构建，快速验证

```
bisheng -x asc test.cpp -o test -I {atvoss_path}/include .....
```

- 采用Host-device融合编译技术
- 减少算子交付件
- 简化编译构建步骤

- 提供一键编译运行的样例执行脚本

```
# 设置环境变量  
source xxx/latest/bin/setenv.bash  
  
# 切换到样例的目录  
cd ../examples  
  
# 编译执行命令  
bash run_examples.sh cast
```

- 样例执行结果：

```
[root@localhost examples]# bash run_examples.sh cast  
Executing with npu mode  
Accuracy verification passed.  
Sample cast passed!
```

# 高效：灵活组装模板实现算子功能

## ⑩ 组装Block模板：

BlockBuilder中嵌套自定义算子Compute

## ⑪ 组装Kernel模板：

KernelBuilder中嵌套BlockOp

## ⑫ 组装Device模板：

DeviceAdapter中嵌套KernelOp

## ⑬ 自定义计算表达式：与模板框架解耦，可独立验证

```
template<typename T>
struct RmsnormCompute {

    template <template <typename> class Tensor>
    __host__ __aicore__ constexpr auto Compute() const
    {
        ...
        return (tmp = x * x,
                z = ReduceSum<Pattern::AR>(tmp),
                z = Divs(z, GetShape<Int<-1>>(x)),
                z = Sqrt(z),
                tmp = Broadcast<Pattern::AB>(z),
                tmp = x / tmp,
                y = g * tmp);
    }
};
```

```
struct RmsnormCompute { ... };

// 组装Block层模板
using BlockOp = EleWise::BlockBuilder<
    Compute      = RmsnormCompute,
    Policy       = blockPolicy,
    ScheduleCfg = EleWise::BlockDefaultCfg,
    Schedule     = EleWise::BlockDefaultSchedule>;


// 组装Kernel层模板
using KernelOp = EleWise::KernelBuilder<
    BlockBuilder = BlockOp,
    Policy       = kernelPolicy,
    ScheduleCfg = EleWise::KernelDefaultCfg,
    Schedule     = EleWise::KernelDefaultSchedule>;

// 组装Device层模板
using DeviceOp = Atvoss::DeviceAdapter<KernelOp>;
```

# 高效：同一个计算表达式，使用不同模板实现多种算子场景

```
// 定义Compute
struct RmsnormCompute {
    template <template <typename> class Tensor>
    __host__ __aicore__ constexpr auto Compute() const
    {
        ...
        return (tmp = x * x,
                z = ReduceSum<Pattern::AR>(tmp),
                z = Divs(z, GetShape<Int<-1>>(x)),
                z = Sqrt(z),
                tmp = Broadcast<Pattern::AB>(z),
                tmp = x / tmp,
                y = g * tmp);
    }
};
```

使用 ATVOSS 前后代码行数对比



\*以上数据来源于华为内置算子，统计核心功能代码行数

## ⑩ 单Tile块切R场景, Reduce轴在Tile块内完成

```
// 组装Block层模板
using BlockOp = EleWise::BlockBuilder<
    Compute      = RmsnormCompute,
    Policy       = blockPolicy,
    ScheduleCfg = EleWise::BlockDefaultCfg,
    Schedule     = EleWise::BlockDefaultSchedule>

// 组装Kernel层模板
using KernelOp = EleWise::KernelBuilder<
    BlockBuilder = BlockOp, ... >;
```

## ⑪ 单核切R场景, Reduce轴在Block内完成

```
// 组装Block层模板
using BlockOp = Reduce::BlockBuilder<
    Compute      = RmsnormCompute,
    Policy       = blockPolicy,
    ScheduleCfg = Reduce::BlockDefaultCfg,
    Schedule     = Reduce::BlockDefaultSchedule>

// 组装Kernel层模板
using KernelOp = EleWise::KernelBuilder<
    BlockBuilder = BlockOp, ... >;
```


# 高性能：

## ■ 模板自动推到最佳切分策略

1. 基于表达式计算存活节点，编译时获取最佳切分
2. 基于最佳切分决定切核策略

```
template<typename T>
struct RmsnormCompute {
    template <template <typename> class Tensor>
    __host__ __aicore__ constexpr auto Compute() const
    {
        ...
        return (tmp = x * x,
                z = ReduceSum<Pattern::AR>(tmp),
                z = Divs(z, GetShape<Int<-1>>(x)),
                z = Sqrt(z),
                tmp = Broadcast<Pattern::AB>(z),
                tmp = x / tmp,
                y = g * tmp);
    }
};
```

## ■ 定制场景，性能达内置算子的 1.2x以上



# 高扩展：参数化配置、可移植

```
struct RmsnormCompute {
    __host__ __aicore__ constexpr auto Compute() const
    {
        ....
        return (tmp = x * x,
                z = ReduceSum<Pattern::AR>(tmp),
                z = Divs(z, GetShape<Int<-1>>(x)),
                z = Sqrt(z),
                tmp = Broadcast<Pattern::AB>(z),
                tmp = x / tmp,
                y = g * tmp);
    }
};
```

## ■ 开放调优参数

支持Tile块大小配置

支持Kernel&block层Policy参数配置

## ■ 表达式模板自由移植

通过Tile层与底层计算实现隔离，不同硬件平台可直接移植

```
// 配置Tile块大小
using TileShape = Atvoss::Shape<256, 1024>;

// 配置Block层静态Policy
static constexpr EleWise::BlockPolicy blockPolicy {
    190 * 1024, // Maximum UB space size.
    TileShape{};
};

// 配置Kernel层静态Policy
static constexpr EleWise::KernelPolicy kernelPolicy {
    48, // 用户自定义最大启动核数
    EleWise::PolicySegment::UNIFORM_SEGMENT // 多核分配策略
};
```

# 高扩展：支持定制模板Schedule与Policy

## ■ Block层模板框架支持高度自定义

```
// 组装Block层模板
using BlockOp = EleWise::BlockBuilder<
    Compute      = AddMulCompute,
    Policy       = blockPolicy,
    ScheduleCfg = EleWise::BlockDefaultCfg,
    Schedule     = EleWise::BlockDefaultSchedule>;
```

- **Compute**: 自定义计算表达式
- **Policy**: 自定义切分策略
- **Schedule& ScheduleCfg** : 定制Block单核处理流程;

## ■ Kernel层模板框架支持高度自定义

```
// 组装Kernel层模板
using KernelOp = EleWise::KernelBuilder<
    BlockBuilder      = BlockOp,
    Policy            = kernelPolicy,
    ScheduleCfg      = EleWise::KernelDefaultCfg,
    Schedule          = EleWise::KernelDefaultSchedule>;
```

- **BlockBuilder** : 自定义计算表达式
- **Policy**: 自定义切分策略
- **Schedule& ScheduleCfg** : 定制Kernel多核处理流程;

# ATVOSS总体架构

五层模块化构建，各级模板可复用、可替换、可定制



# ATVOSS编程界面

```
template<typename T>
struct AddMulCompute {

    template <template <typename> class Tensor>
    __host__ __aicore__ constexpr auto Compute() const
    {
        ...
        return (tmp = x * x, ...);
    }
};

// 组装模板
using BlockOp = EleWise::BlockBuilder<AddMulCompute, blockPolicy, ...>;
using KernelOp = EleWise::KernelBuilder<BlockOp, kernelPolicy, ...>;
using DeviceOp = Atvooss::DeviceAdapter<KernelOp>;

int main() {

    // 构造输入数据, 示例是固定输入数据值的, 实际可自定义生成数据
    Atvooss::Tensor<float> t1(dataV1/*data host ptr*/, {32, 32} /*shape*/);
    Atvooss::Tensor<float> t2(dataV2/*data host ptr*/, {32, 32} /*shape*/);
    Atvooss::Tensor<float> t3(dataV3/*data host ptr*/, {32, 32} /*shape*/);

    // 生成入参信息
    Arguments arguments = ArgumentsBuilder{}.input({t1, t2}).output({t3}).build();

    // 算子执行
    DeviceOp addMulOp();
    addMulOp.Run(arguments);
}
```

用户自定义计算表达式, 描述计算逻辑

用户配置Tile块大小、 UB最大内存等

用户自定义构造host侧输入输出数据

CANN

# 目录

Part 1 ATVOSS

Part 2 ATVOSS开源仓

# ATVOSS仓介绍

```
└── build.sh  
└── cmake  
└── CMakeLists.txt  
└── CONTRIBUTING.md  
└── docs  
└── examples  
└── include  
    ├── device  
    ├── kernel  
    ├── block  
    ├── tile  
    └── utils  
└── LICENSE  
└── OAT.xml  
└── README.md  
└── scripts  
└── SECURITY.md  
└── test  
└── Third_Party_Open_Source_Software_Note  
└── version.info
```



- **Docs**: 包含快速上手、用户使用指南等文档
- **Examples**: 算子开发样例
- **include**: 模板实现头文件库

包含各层模板实现: device/kernel/block/tile

包含计算表达模板等基础工具实现: utils

...

代码仓链接: <https://gitcode.com/cann/atvoss>

# ATVOSS未来规划: 更丰富、更优、更易用



# 交流 & 贡献



ATVOSS 代码仓  
<https://gitcode.com/cann/atvooss>

欢迎交流！欢迎贡献！欢迎共创合作！

# Thank you.

社区愿景：打造开放易用、技术领先的AI算力新生态

社区使命：使能开发者基于CANN社区自主研究创新，构筑根深叶茂、跨产业协同共享共赢的CANN生态

Vision: Building an Open, Easy-to-Use, and Technology-leading AI Computing Ecosystem  
Mission: Enable developers to independently research and innovate based on the CANN community and build a win-win CANN ecosystem with deep roots and cross-industry collaboration and sharing.



上CANN社区获取干货



关注CANN公众号获取资讯

<https://gitcode.com/cann>

**CANN**