

TileLang Ascend开发入门

<https://gitcode.com/cann>

CANN

目录

Part 1 TileLang介绍

Part 2 TileLang Ascend介绍

Part 3 TileLang Ascend开发入门

Part 4 TileLang Ascend 后续计划

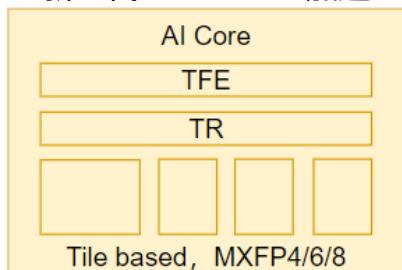
趋势与机遇：基于Tile的软硬件协同

AI加速器发展的新趋势使得基于Tile编程语言变得越来越重要。

Tile-AI加速硬件

- 更高的数据传输和计算效率
- 本地大容量SRAM，再结合3D DRAM，片上互联，提供超高计算带宽
- 要求以更大的Tile粒度来进行指令集设计和体系结构设计

新一代Tile-based AI加速



片上网络

Tile-编程语言

- 在更大的Tile粒度上进行编程和编译，可以更方便和高效地优化模型算子
- 实现Tile编程计算抽象和Tile硬件抽象一致，高效协同

Tile Programming

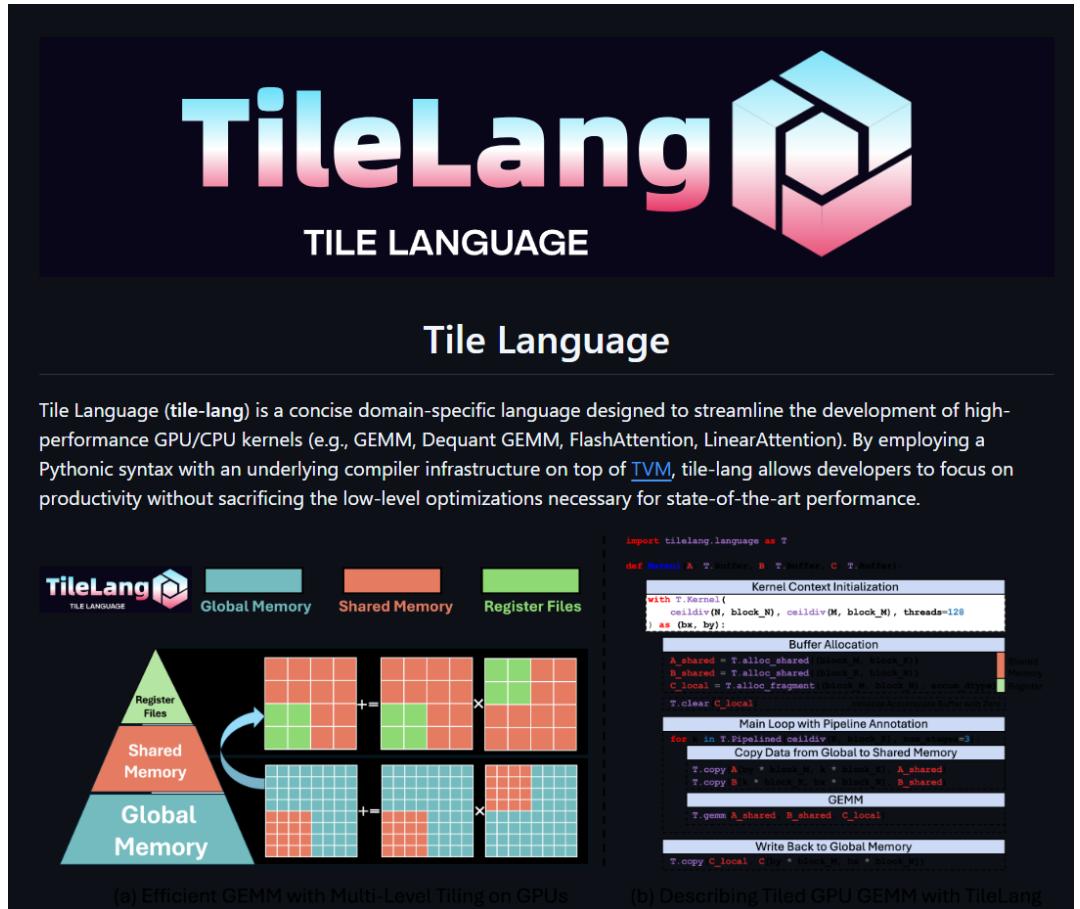
Block-wide cooperative execution on regular Tiles of data

Data & operation granularity is array / tensor



Tile编程语言： TileLang

- Tile-level的类Python的AI编程语言 (DSL)
 - 开发更简单
 - 提供了更好的性能
 - 可以更加细粒度地控制计算调度
- 面对 Tile 硬件趋势的兴起，推出Tile编程语言TileLang，主要核心技术团队来自北京大学
- 2025-01-20开源，现已获得超过3.7K stars
<https://github.com/tile-ai/tilelang>
- 支持多硬件后端
GPU、NPU、TPU、CPUs

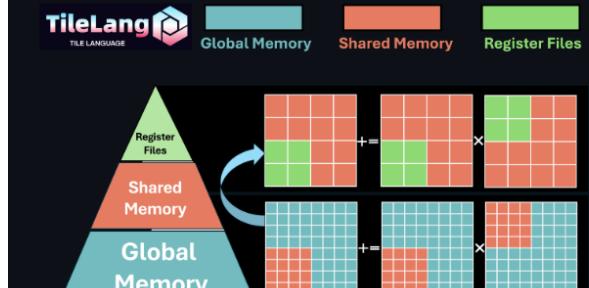


The image shows the TileLang logo at the top, followed by a detailed description of the Tile Language. It includes a diagram illustrating multi-level tiling for GEMM operations across Global Memory, Shared Memory, and Register Files. Below this, two code snippets demonstrate how to describe tiled GPU GEMM operations using TileLang.

Tile Language

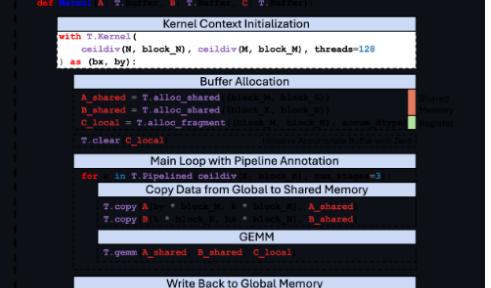
Tile Language (tile-lang) is a concise domain-specific language designed to streamline the development of high-performance GPU/CPU kernels (e.g., GEMM, Dequant GEMM, FlashAttention, LinearAttention). By employing a Pythonic syntax with an underlying compiler infrastructure on top of [TVM](#), tile-lang allows developers to focus on productivity without sacrificing the low-level optimizations necessary for state-of-the-art performance.

(a) Efficient GEMM with Multi-Level Tiling on GPUs



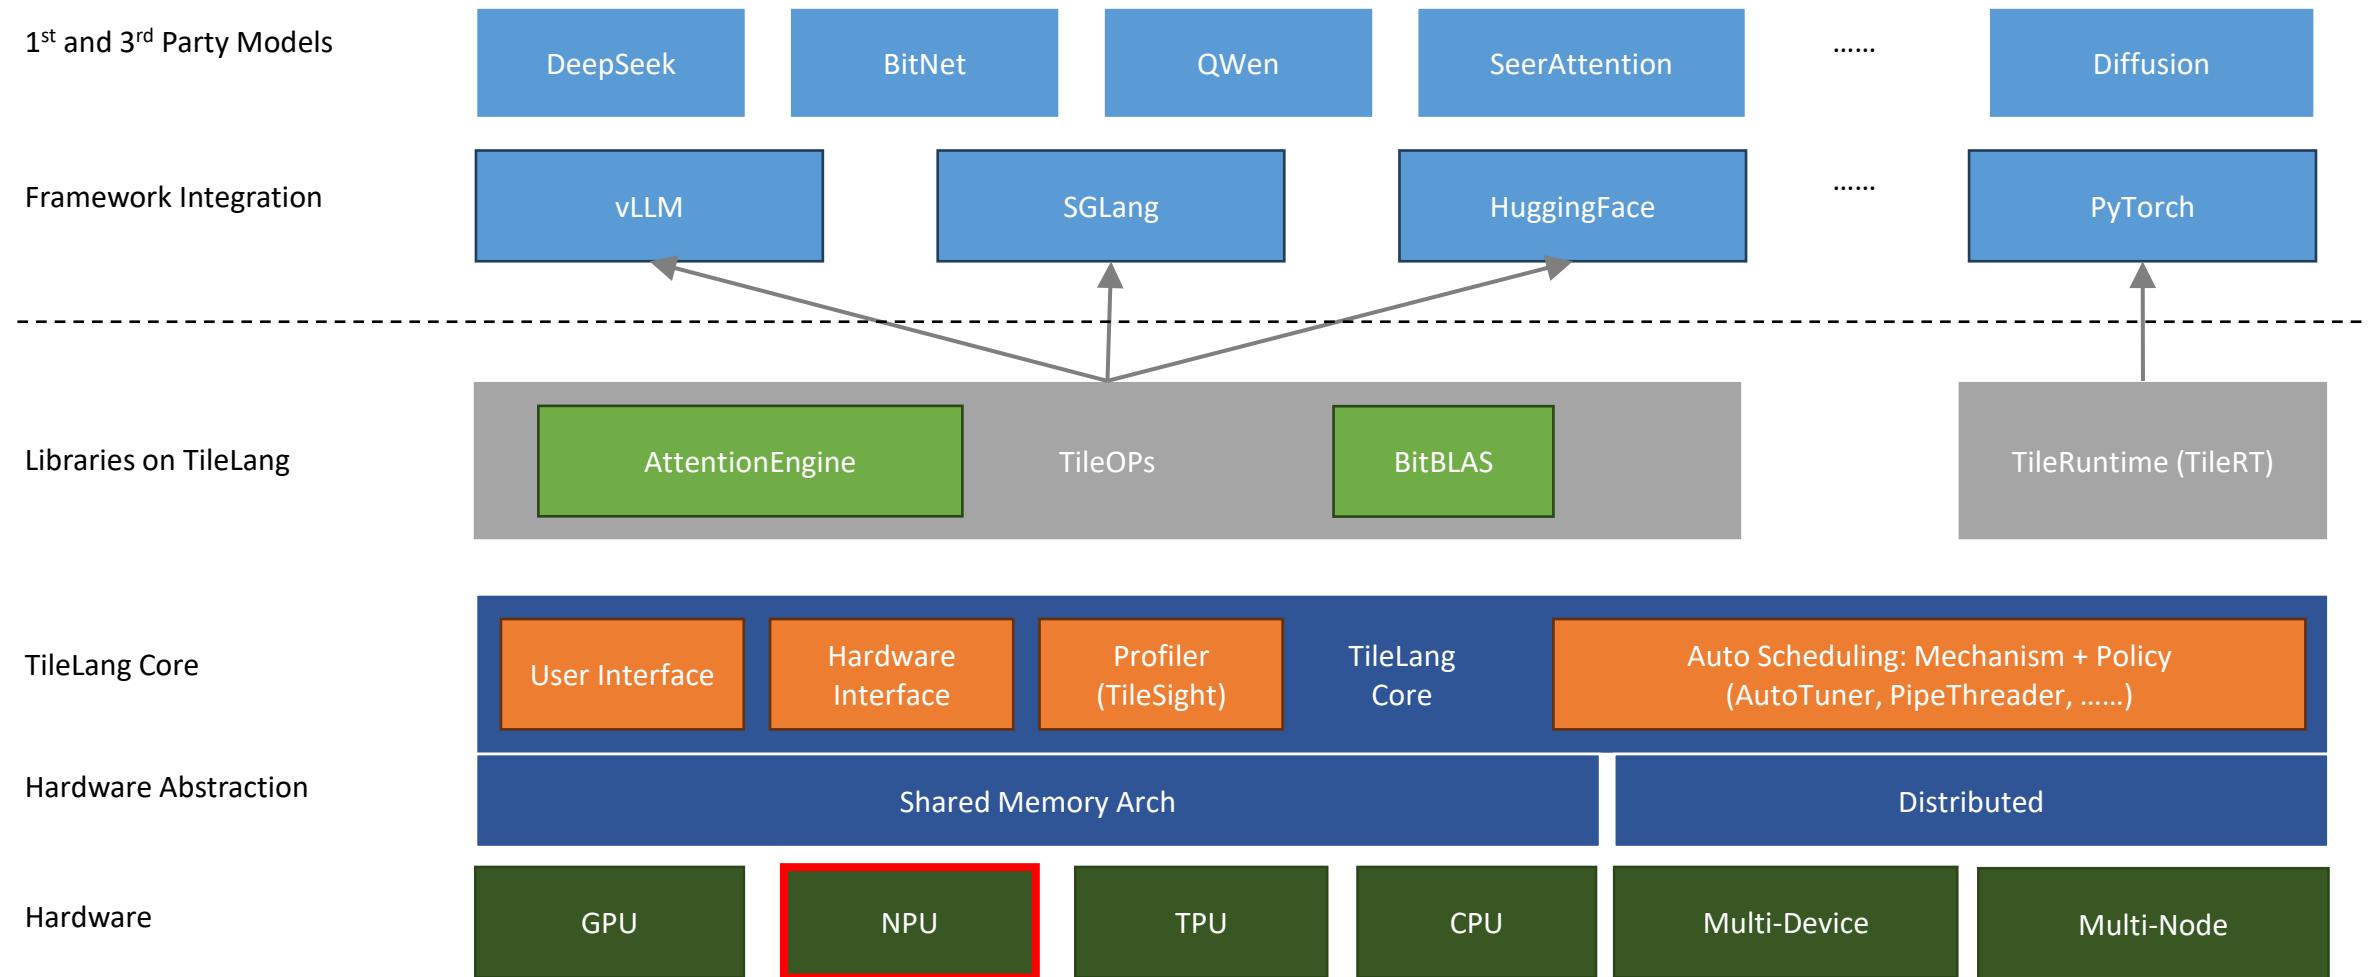
The diagram illustrates a GEMM operation with three input matrices A, B, and C. Matrix A is loaded from Global Memory into Shared Memory. Matrix B is also loaded from Global Memory into Shared Memory. Matrix C is initially in Shared Memory and is updated via a pipeline. The diagram shows the flow of data between Global Memory, Shared Memory, and Register Files, with annotations for addition (+=) and multiplication (x).

(b) Describing Tiled GPU GEMM with TileLang



```
import tilelang.language as T
def Matmul A T.buffer B T.buffer C T.buffer:
    Kernel Context Initialization
    with T.Kernel(
        ceildiv(N, block_N), ceildiv(M, block_M), threads=128
    ) as (tx, by):
        Buffer Allocation
        A_shared = T.alloch_shared(tx * block_N, tx * block_M)
        B_shared = T.alloch_shared(tx * block_N, by * block_M)
        C_local = T.alloch_fragment(tx * block_N, by * block_M, access_dtypes=[float])
        T.clear C_local
        Main Loop with Pipeline Annotation
        for int T.Pipelined ceildiv(tx * block_N, block_N), num_stages=3:
            Copy Data from Global to Shared Memory
            T.copy A A_shared
            T.copy B B_shared
            GEMM
            T.gemm A_shared B_shared C_local
            Write Back to Global Memory
            T.copy C_local C[tx * block_N, by * block_M]
```

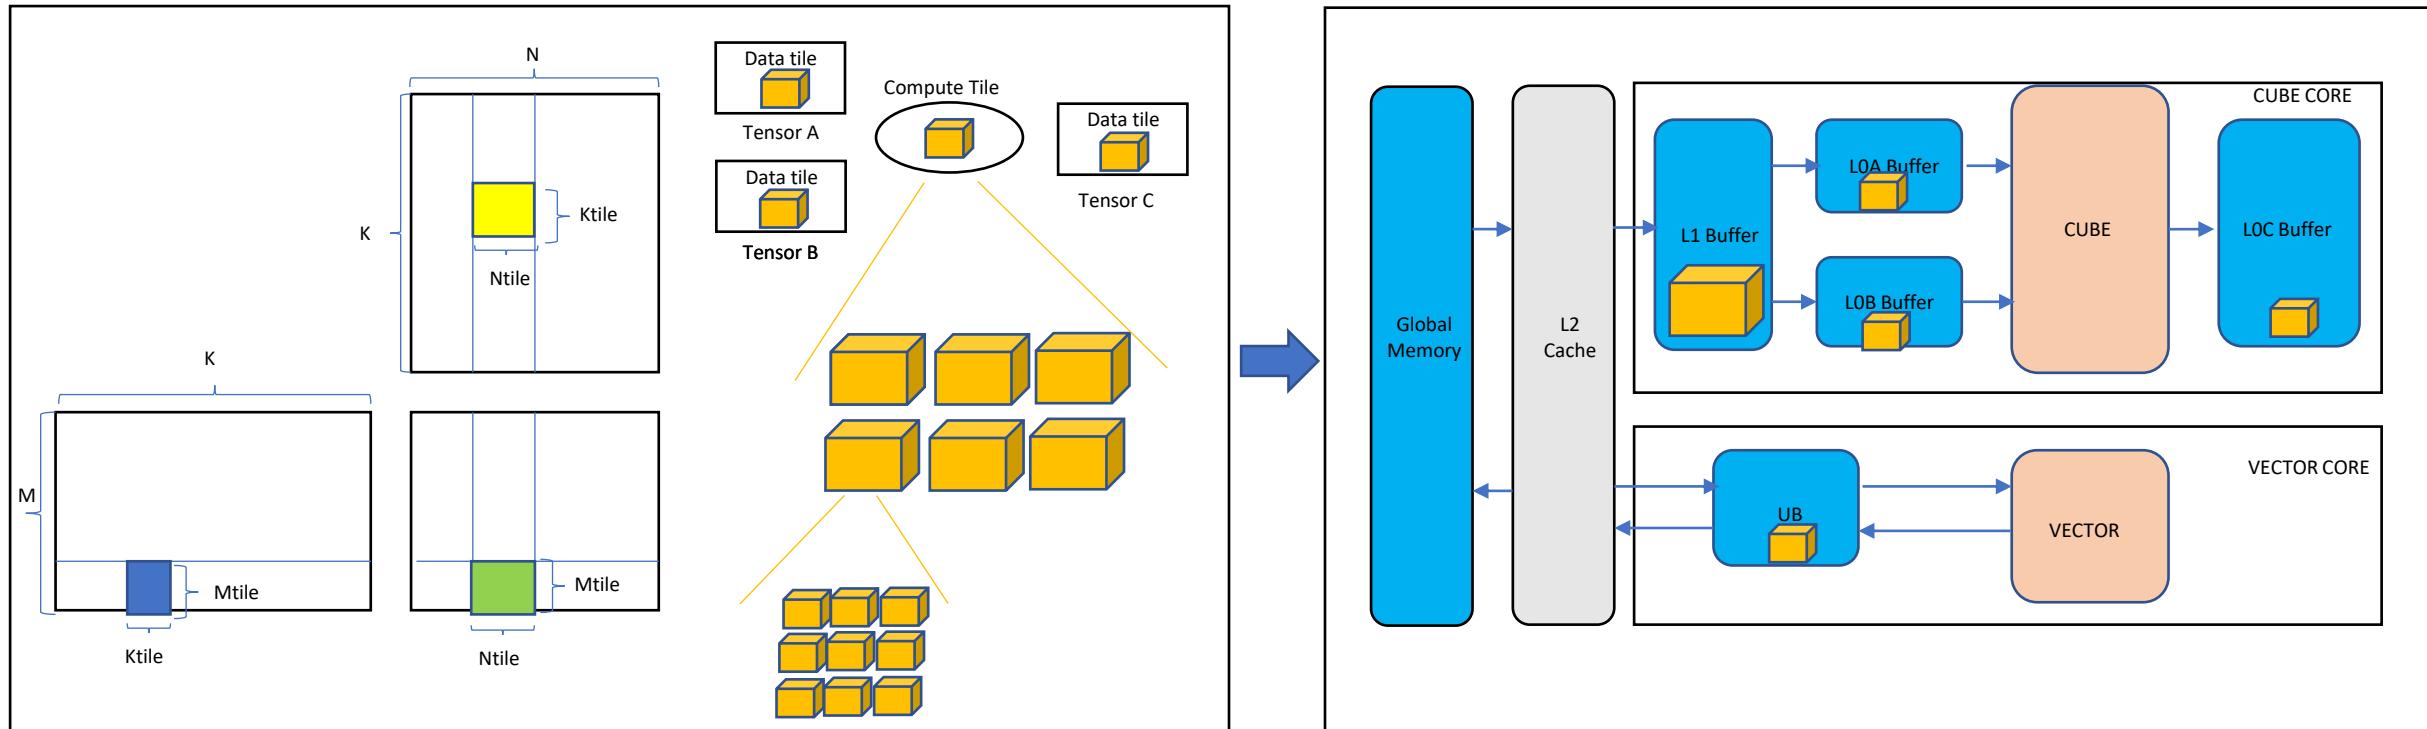
TileLang系统和生态



TileLang：核心设计理念

TileLang = Tile-VM + API原语

Tile-VM = Tile Task 模型 + Tile 硬件抽象

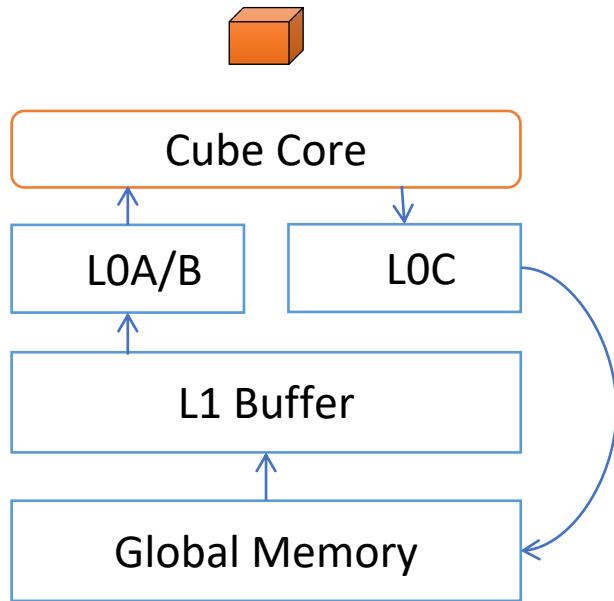


层次化 Tile Task模型

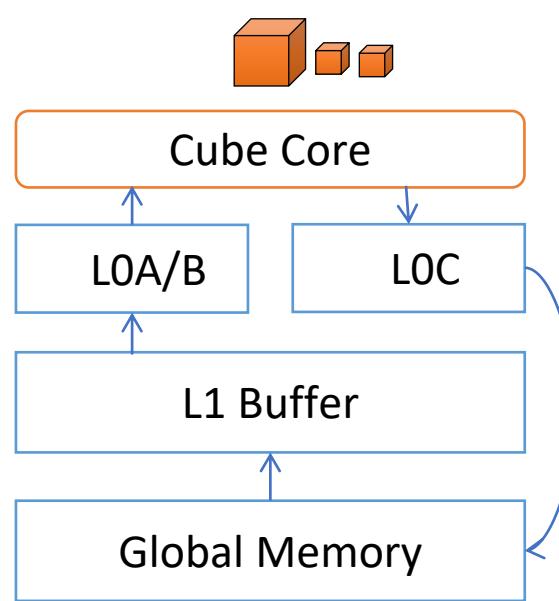
Tile 硬件抽象

TileLang API原语： Tile Allocation

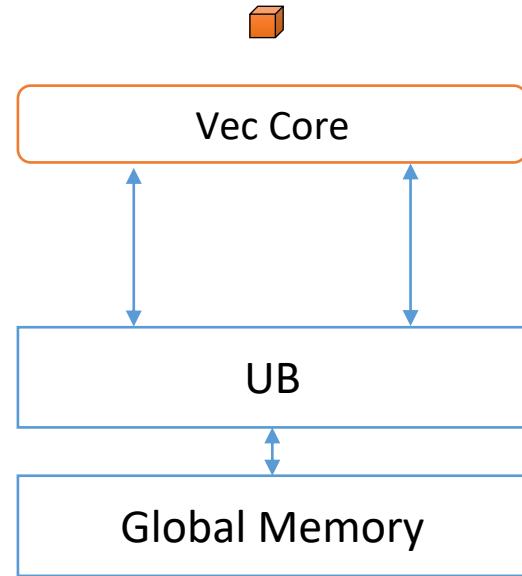
- T.alloc_L1



- T.alloc_l0a/b/c

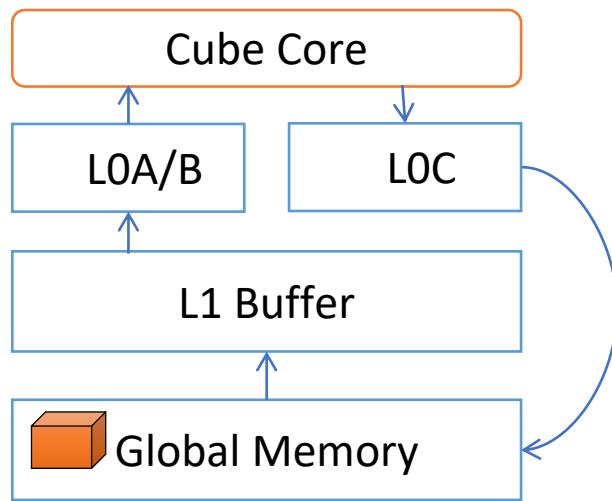


- T.alloc_UB

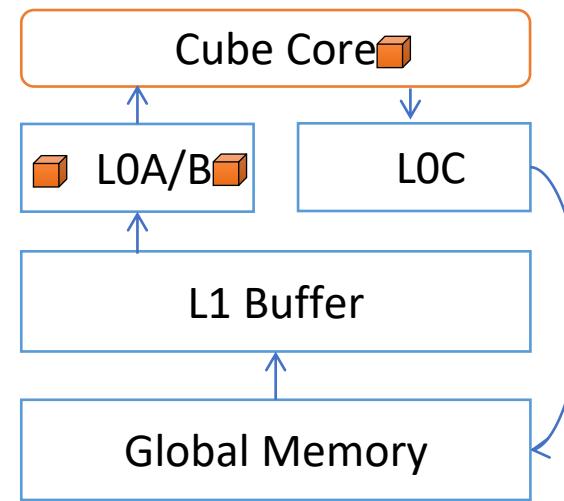


TileLang API原语： Tile Operation

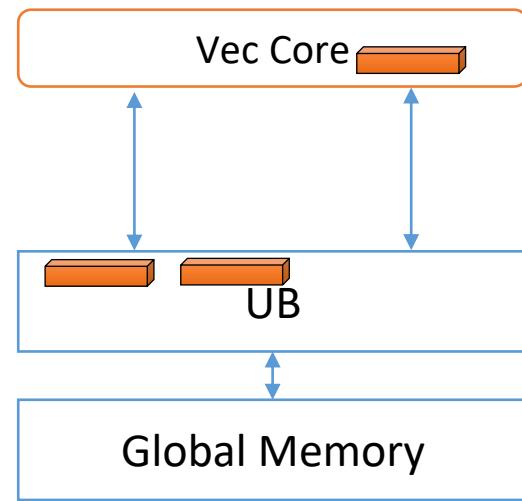
- T.copy()



- T.gemm()



- T.reduce()



目录

Part 1 TileLang介绍

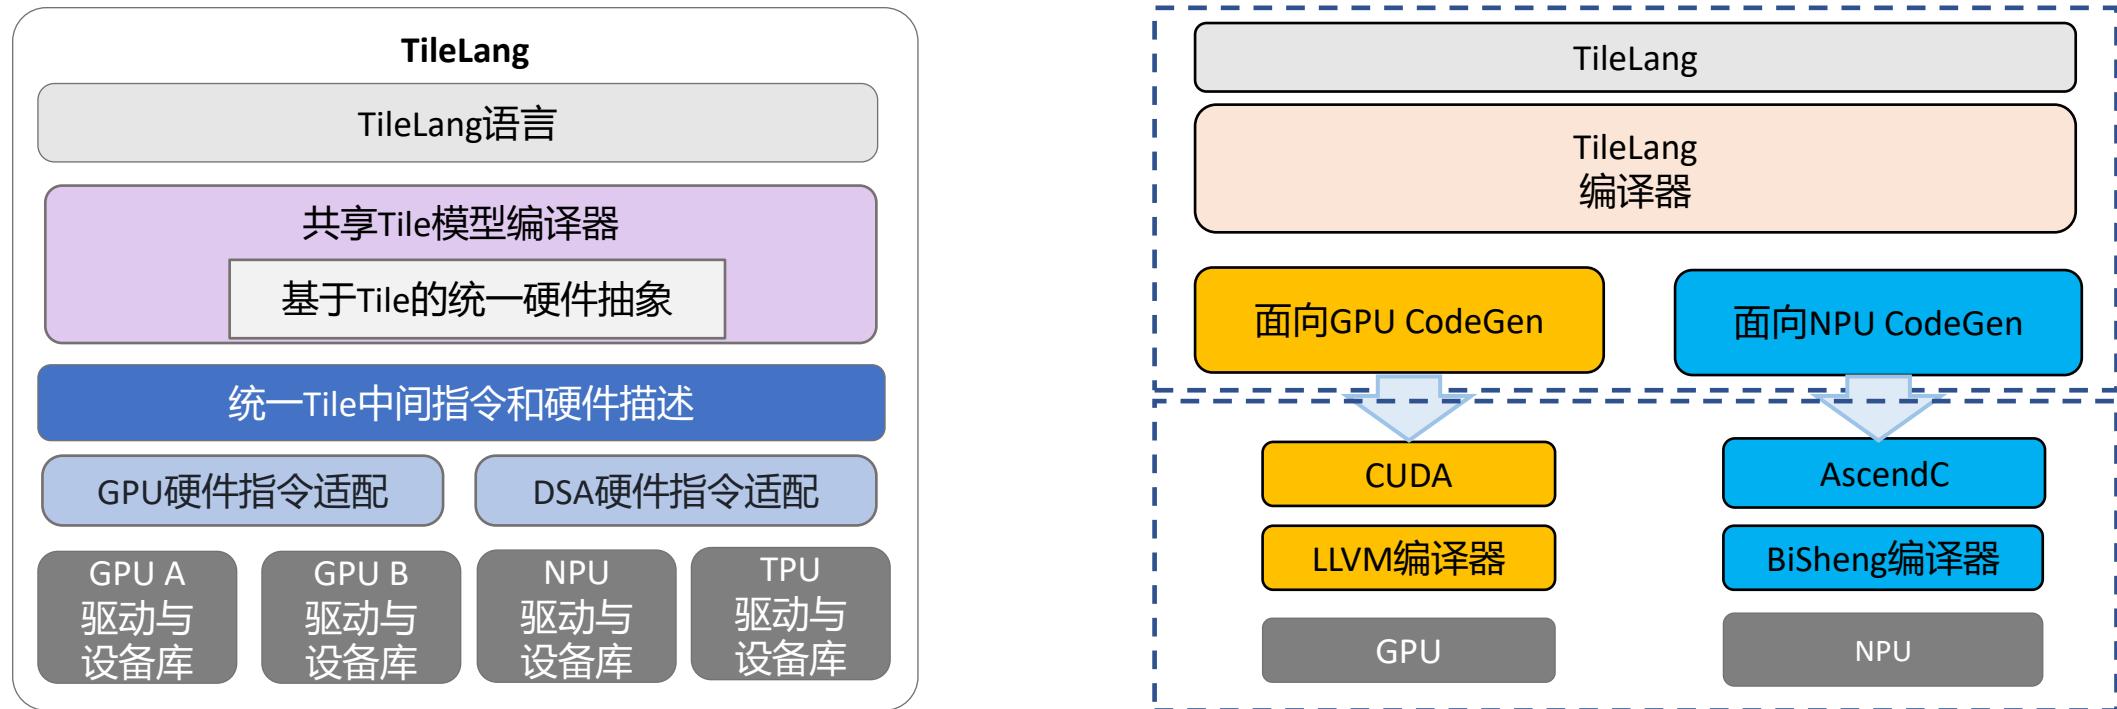
Part 2 TileLang Ascend介绍

Part 3 TileLang Ascend开发入门

Part 4 TileLang Ascend 后续计划

TileLang Ascend适配

Ascend适配TileLang：基于Tile的统一硬件表达的抽象，提供轻量级的底层装换，并提供扩展接口提供Ascend专属能力。



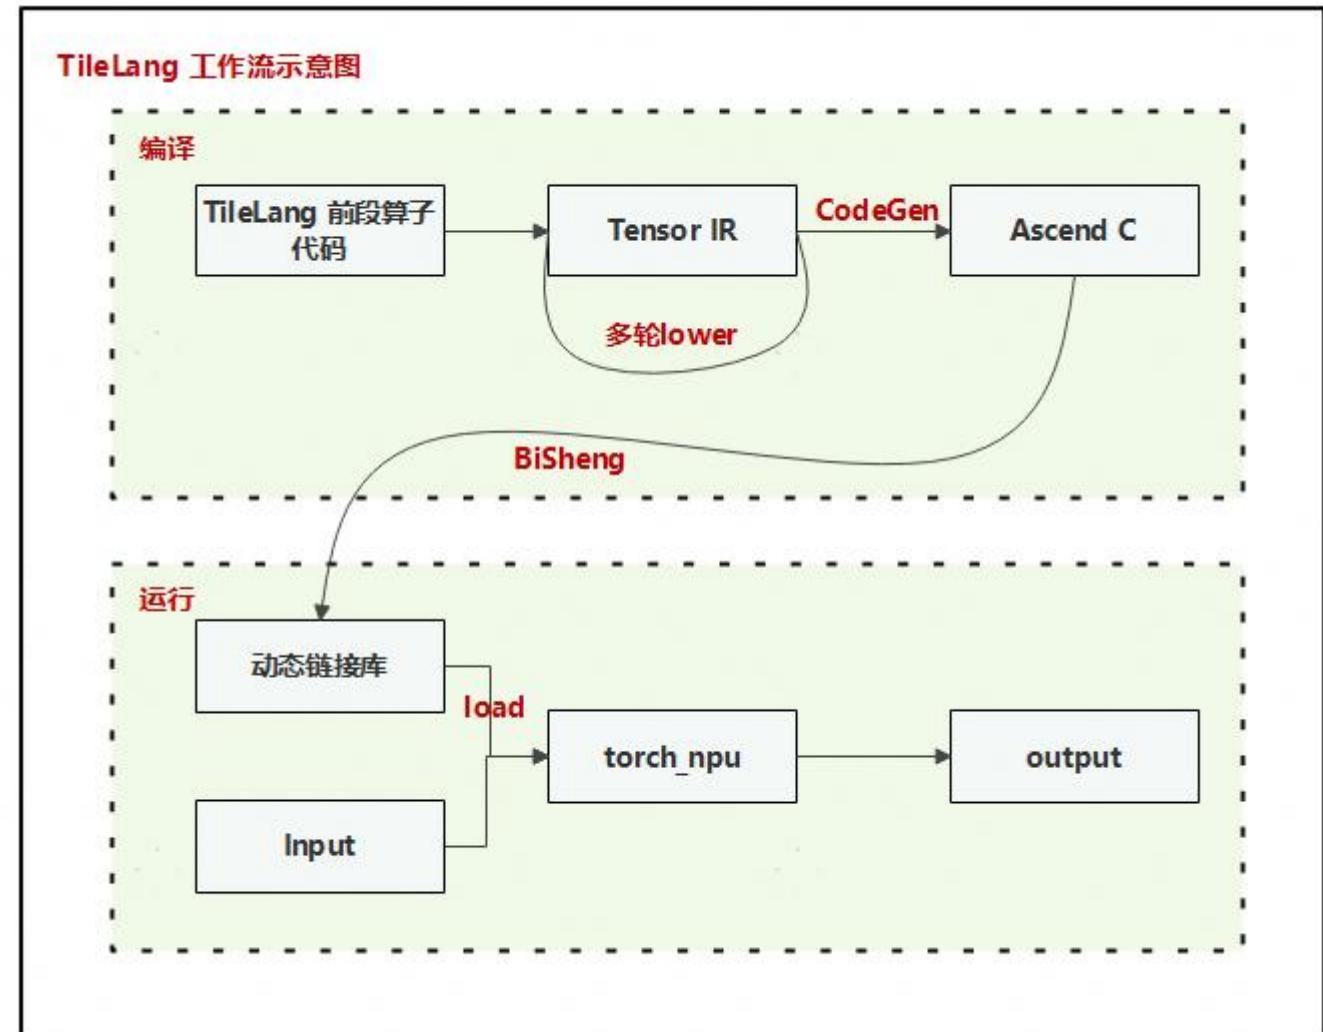
TileLang Ascend工作流程

编译阶段

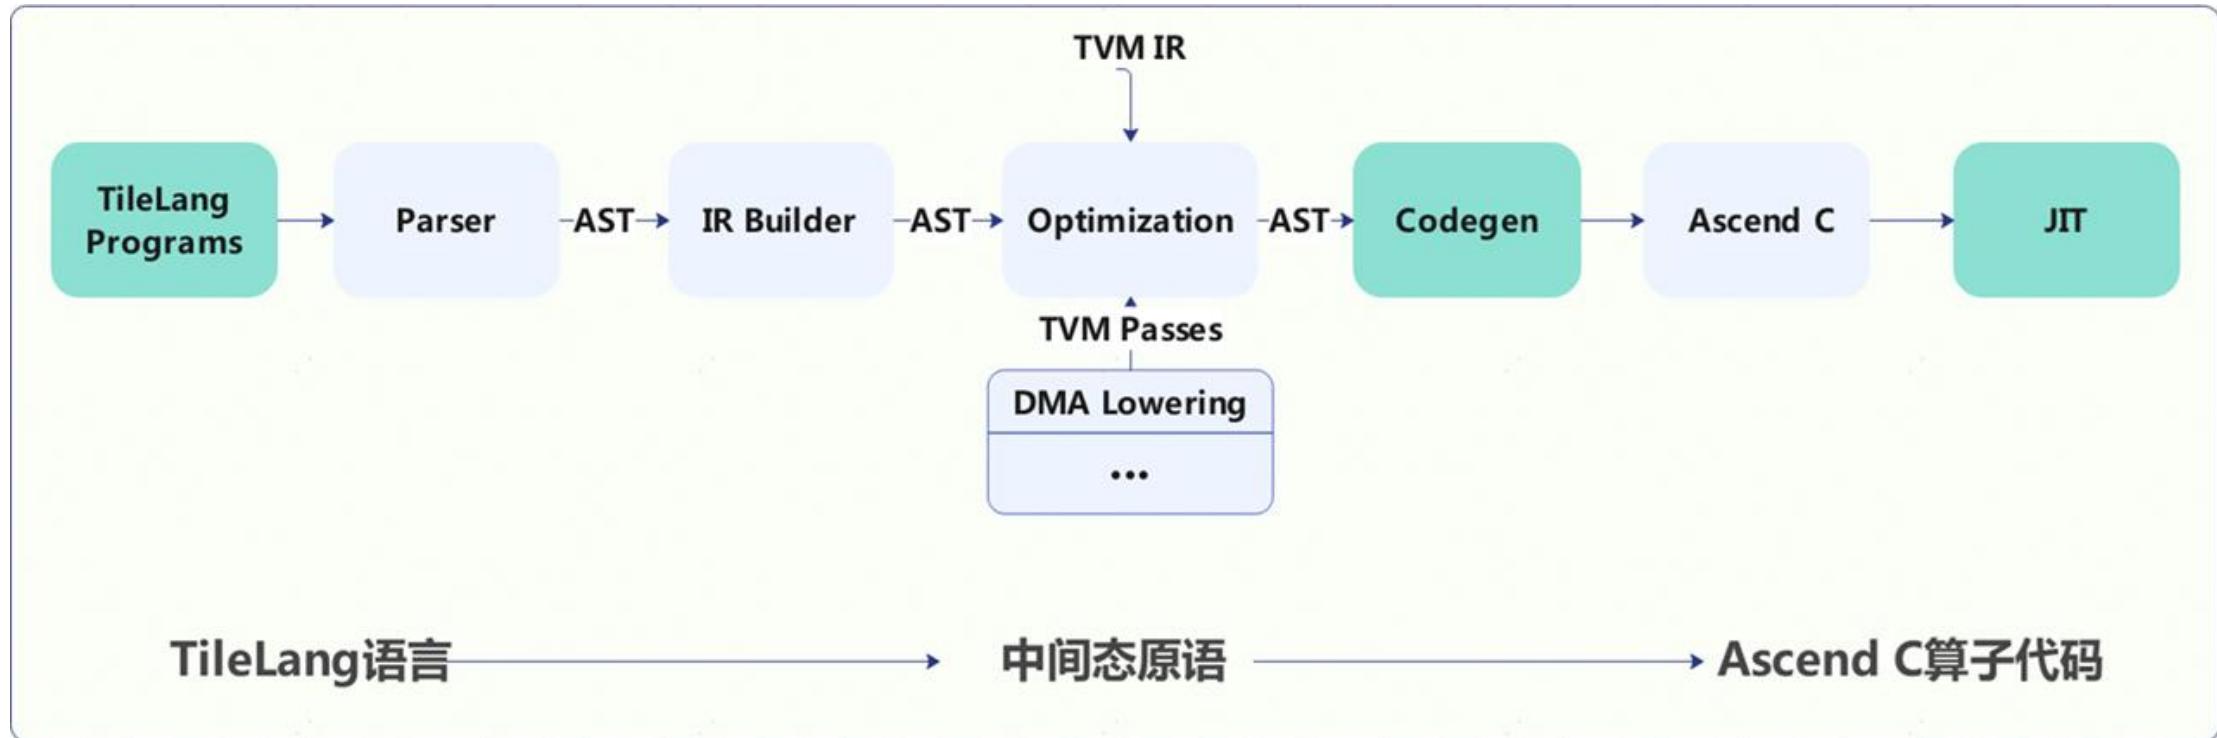
- 多级Lowering转换：TileLang算子根据NPU硬件特性进行多级降级，生成针对昇腾硬件优化的TensorIR表示。
- AscendC代码生成：基于TensorIR，使用专门的Ascend CodeGen模块生成对应的AscendC代码。
- 动态库编译：通过毕昇编译器（BiSheng）将AscendC代码编译成动态链接库（.so文件）。

运行阶段

- 库文件加载：通过torch_npu运行时库将.so文件加载到Python环境中
- 函数封装：将算子封装为可调用的Python函数对象
- 执行调用：用户以普通Python函数方式调用，提供输入张量即可在昇腾NPU上执行计算并获得结果



TileLang Ascend编译运行流程



- **CodeGen**

通过预定义的AscendC指令模板，封装NPU底层硬件操作（如AI Core计算指令、DMAC数据传输指令）的语法细节与参数约束，CodeGen可自动解析TileLang语义树，根据语义逻辑匹配对应的指令模板，并通过指令映射规则，实现TileLang与AscendC NPU后端的无缝适配。

- **JIT**

通过JIT调用CodeGen生成AscendC代码，并对整个过程进行动态调控，解决静态编译的局限性，确保生成的AscendC代码动态适配NPU特性，同时最大化提升算子执行效率。

TileLang Ascend原语介绍

原语类型	原语名称	用法示例	功能介绍
内核原语	kernel	T.kernel(block_num, is_npu=True) as (cid, vid)	该原语对应到Ascend C的kernel调用处，其中block_num对应于<<< >>>中的第一个参数，表示这个kernel开启多少个子任务数量。cid的范围为[0,block_num)[0,block_num)[0,block_num), vid的范围为0或1。因为A2的cv核默认配比为1:2, 可以通过vid指定当前vector的索引。
内存分配原语	alloc_L1/L0A/L0B/L0C/UB	T.alloc_L1/L0A/L0B/L0C/UB(shape, dtype)	用于分配位于片上内存的buffer；通过指定shape和数据类型标记buffer的信息。
数据搬运原语	copy	T.copy(src, dst)	将src上的buffer拷贝到dst上，注意buffer可以通过BufferLoad或者BufferRegion指定一小块区域。
计算原语	gemm, add, mul, reduce_max...	T.reduce_max(dst, src, tmp, dim)	其中dim指定为对应规约的维度，目前只支持二维的规约。
同步原语	set_flag, wait_flag, set_cross_flag, wait_cross_flag, pipe_barrier	T.set_cross_flag(pipe: str, eventId: int)	其中pipe为需要同步的流水线，eventId为同步事件编号。

目录

Part 1 TileLang介绍

Part 2 TileLang Ascend介绍

Part 3 TileLang Ascend开发入门

Part 4 TileLang Ascend 后续计划

TileLang基础语法结构

```
@tilelang.jit(out_idx=[-1])
def matmul(M, N, K, block_M, block_N, K_L1, dtype="float16", accum_dtype="float"):
    m_num = M // block_M
    n_num = N // block_N

    @T.prim_func
    def main():
        A: T.Tensor((M, K), dtype)
        B: T.Tensor((K, N), dtype),
        C: T.Tensor((M, N), dtype),
    ):
        with T.Kernel(m_num * n_num, is_npu=True) as (cid, ):
            bx = cid // n_num
            by = cid % n_num

            A_L1 = T.alloc_L1((block_M, K_L1), dtype)
            B_L1 = T.alloc_L1((K_L1, block_N), dtype)

            C_L0 = T.alloc_L0C((block_M, block_N), accum_dtype)

            with T.Scope("C"):
                loop_k = T.ceildiv(K, K_L1)
                for k in t.serial(loop_k):
                    T.copy(A[bx * block_M, k * K_L1], A_L1)
                    T.copy(B[k * K_L1, by * block_N], B_L1)

                    T.barrier_all()
                    T.gemm_v0(A_L1, B_L1, C_L0, init=(k == 0))

                    T.barrier_all()

                    T.copy(C_L0, C[bx * block_M, by * block_N])

    return main

func = matmul(M, N, K, 128, 256, 64)
a = torch.randn(M, K).half().npu()
b = torch.randn(K, N).half().npu()
c = torch.empty(M, N).half().npu()
c = func(a, b)
```

@tilelang.jit 装饰器标识函数会被jit动态编译，out_idx表示主计算函数返回参数索引。

@T.prim_func 装饰器在模块内定义主计算函数。

T.Tensor 声明数据类型Tensor数据缓冲区，并指定其形状和数据类型。

T.Kernel原语触发算子kernel调用，对应AscendC的kernel调用。

T.Scope标识代码的计算单元，“C” 表示在cube核运行，“V” 表示在vector核上运行。

算子具体计算逻辑。

返回主计算函数对象供调用。

在主程序中调用TileLang算子，触发jit编译。
调用主计算函数。

CANN

TileLang样例：GEMM_Basic

```
@tilelang.jit(out_idx=[-1])
def matmul(M, N, K, block_M, block_N, K_L1, dtype="float16", accum_dtype="float"):
    m_num = M // block_M
    n_num = N // block_N

    @T.prim_func
    def main(
        A: T.Tensor((M, K), dtype),
        B: T.Tensor((K, N), dtype),
        C: T.Tensor((M, N), dtype),
    ):
        with T.Kernel(m_num * n_num, is_npu=True) as (cid, _):
            bx = cid // n_num # 计算BlockM偏移量
            by = cid % n_num # 计算BlockN偏移量

            A_L1 = T.alloc_L1((block_M, K_L1), dtype) # 申请L1空间
            B_L1 = T.alloc_L1((K_L1, block_N), dtype)

            C_L0 = T.alloc_L0C((block_M, block_N), accum_dtype) # 申请L0C空间

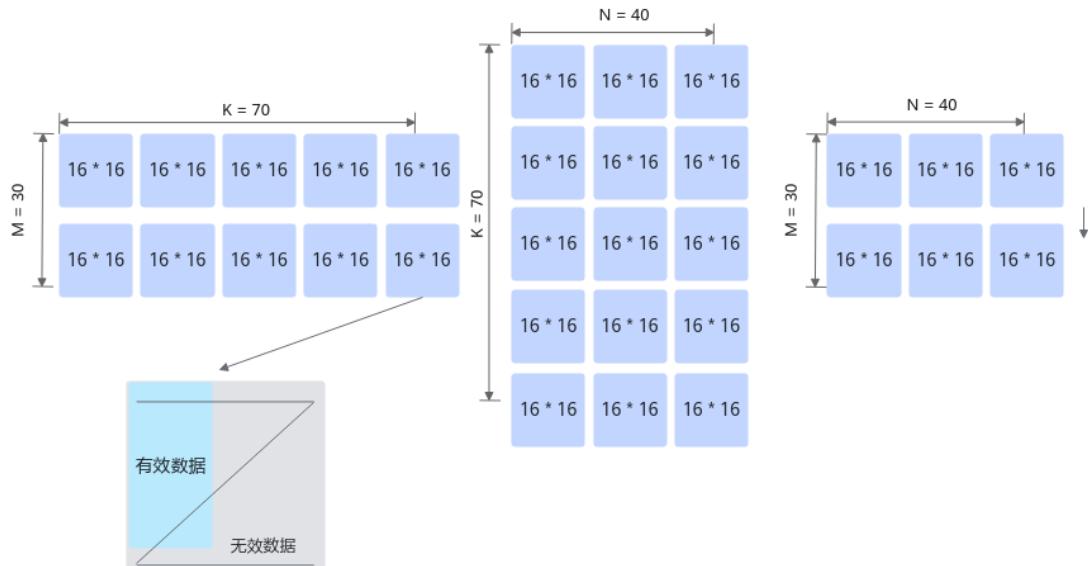
            with T.Scope("C"):
                loop_k = T.ceildiv(K, K_L1) # 计算k轴上按基本块大小K_L1切分的次数
                for k in T.serial(loop_k):
                    T.copy(A[bx * block_M, k * K_L1], A_L1)
                    T.copy(B[k * K_L1, by * block_N], B_L1)

                    T.barrier_all()
                    T.gemm_v0(A_L1, B_L1, C_L0, init=(k == 0)) # 当k为0时表示第一块基本块计算，将
                    LOC清0
                    T.barrier_all()

                    T.copy(C_L0, C[bx * block_M, by * block_N])

    return main
```

左侧为tilelang实现的matmul算子的基础样例。
主要功能：
实现A和B的分块矩阵乘法： $C = A @ B$ 。



TileLang样例： GEMM_Highperf

```

@tilelang.jit(out_idx=[-1]) # 此装饰器会被jit动态编译, out_idx表示kernel返回值索引
def matmul(M, N, K, block_M, block_N, block_K, K_L1, S1, S2, dtype="float16", accum_dtype="float"):
    m_num = M // block_M # M和N方向切分基本块数目
    n_num = N // block_N

    core_num = 20

    @T.macro
    def init_flag(): # 同步相关
        T.set_flag("mte1", "mte2", 0)
        T.set_flag("mte1", "mte2", 1)
        T.set_flag("m", "mte1", 0)
        T.set_flag("m", "mte1", 1)
        T.set_flag("fix", "m", 0)

    @T.macro
    def clear_flag():
        T.wait_flag("mte1", "mte2", 0)
        T.wait_flag("mte1", "mte2", 1)
        T.wait_flag("m", "mte1", 0)
        T.wait_flag("m", "mte1", 1)
        T.wait_flag("fix", "m", 0)

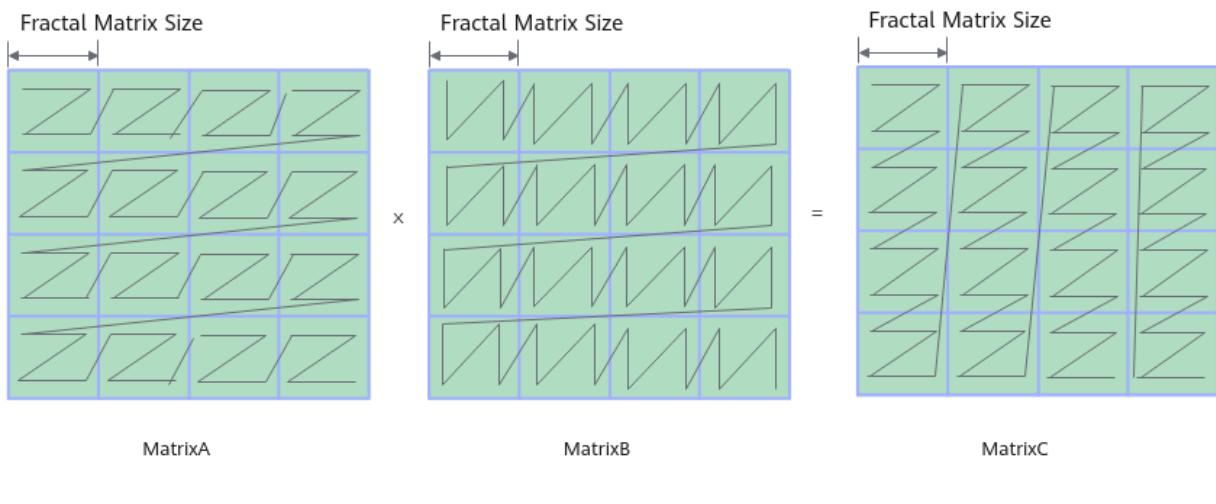
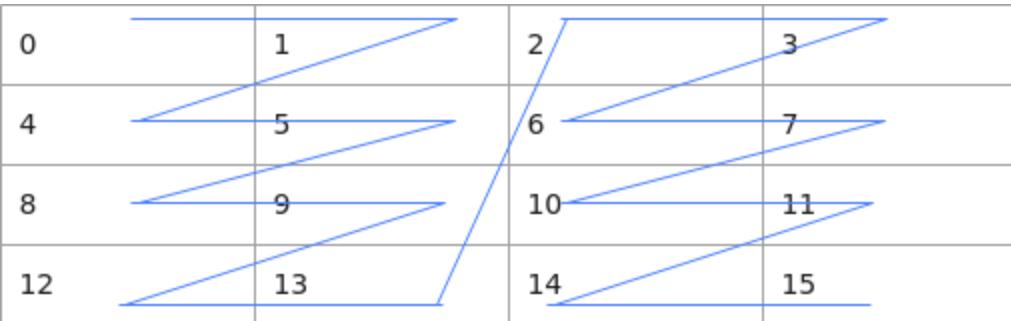
    @T.prim_func
    def main(
        A: T.Tensor((M, K), dtype),
        B: T.Tensor((K, N), dtype),
        C: T.Tensor((M, N), dtype),
    ):
        with T.Kernel(core_num, is_npu=True) as (cid, _):
            A_L1 = T.alloc_L1((S1, block_M, K_L1), dtype) # L1空间申请, S1表示L1上DB数目
            B_L1 = T.alloc_L1((S1, K_L1, block_N), dtype)

            T.annotate_layout({# 指定L1上的内存排布, 后续将tilelang前端的偏移量自动转化为实际内存的偏移量
                A_L1: make_zn_layout(A_L1),
                B_L1: make_zn_layout(B_L1),
            })

            A_L0 = T.alloc_L0A((S2, block_M, block_K), dtype) # L0空间申请, S2表示L2上DB数目
            B_L0 = T.alloc_L0B((S2, block_K, block_N), dtype)
            C_L0 = T.alloc_L0C((block_M, block_N), accum_dtype)

```

Nd->Nz示意：



CANN

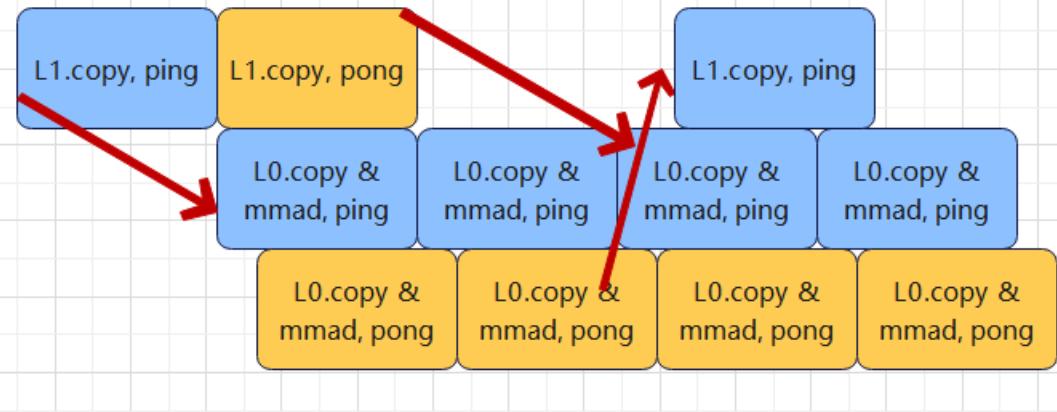
TileLang样例： GEMM_Highperf

```
with T.Scope("C"):
    init_flag() # 发射初始同步
    for i in T.serial(T.ceildiv(m_num * n_num, core_num)):
        T.use_swizzle(i * core_num + cid, M, N, K, block_M, block_N, off=3, in_loop=True) # 增加L2cache命中率
        bx = cid // n_num # 计算BlockM偏移量
        by = cid % n_num # 计算BlockN偏移量
        loop_k = T.ceildiv(K, K_L1) # 计算k方向在l1上的循环次数
        T.wait_flag("mte1", "mte2", 0)
        T.copy(A[bx * block_M, 0], A_L1[0, :, :]) # 提前搬运矩阵基本块，先使用第一块L1buffer
        T.copy(B[0, by * block_N], B_L1[0, :, :])
        T.set_flag("mte2", "mte1", 0)
        T.wait_flag("fix", "m", 0)
        for k in T.serial(loop_k): # 计算k方向在l1上的循环次数
            if k < loop_k - 1: # 由于提前搬运了一次，不执行完毕
                T.wait_flag("mte1", "mte2", (k + 1) % S1)
                T.copy(A[bx * block_M, (k + 1) * K_L1], A_L1[(k + 1) % S1, :, :]) # 按照 pingpong进行搬运
                T.copy(B[(k + 1) * K_L1, by * block_N], B_L1[(k + 1) % S1, :, :])
                T.set_flag("mte2", "mte1", (k + 1) % S1)
            loop_kk = T.ceildiv(K_L1, block_K) # 计算L0上循环次数
            for kk in T.serial(loop_kk):
                if kk == 0: # 该分支用来等待提前发射的同步
                    T.wait_flag("mte2", "mte1", k % S1)
                    T.wait_flag("m", "mte1", kk % S2)
                T.copy(A_L1[k % S1, 0, kk * block_K], A_L0[kk % S2, :, :]) # 计算L0上循环次数
                T.copy(B_L1[k % S1, kk * block_K, 0], B_L0[kk % S2, :, :])
                if kk == 3: # L0全部搬运完毕，通知L1继续搬运
                    T.set_flag("mte1", "mte2", k % S1)
                    T.set_flag("mte1", "m", kk % S2)
                    T.wait_flag("mte1", "m", kk % S2)
                if k == 0 and kk == 0: # 此分支表示LOC清0， else分支表示LOC累加
                    T.mma(A_L0[kk % S2, :, :], B_L0[kk % S2, :, :], C_L0, init=True)
                else:
                    T.mma(A_L0[kk % S2, :, :], B_L0[kk % S2, :, :], C_L0)
                T.set_flag("m", "mte1", kk % S2)

        T.set_flag("m", "fix", 0)
        T.wait_flag("m", "fix", 0)
        T.copy(C_L0, C[bx * block_M, by * block_N]) # 计算完毕，根据偏移量搬运到GM中对应位置
        T.set_flag("fix", "m", 0)
        clear_flag()
        T.barrier_all()
```

下方为左侧代码流水示意图：

流水示意：以L1_BLK=256,L0_BLK=64为例



CANN

TileLang样例：GEMM_函数调用

```
func = matmul(M, N, K, 128, 256, 64, 256, 2, 2)
torch.manual_seed(0)

a = torch.randn(M, K).half().npu()
b = torch.randn(K, N).half().npu()

c = func(a, b)
ref_c = a @ b

torch.npu.synchronize()

torch.testing.assert_close(c, ref_c, rtol=1e-2, atol=1e-2)
print("Kernel Output Match!")

tilelang_time = do_bench(lambda: func(a, b))
torch_time = do_bench(lambda: a @ b)

print(f"tilelang time: {tilelang_time} ms")
print(f"torch time: {torch_time} ms")
```

与torch_npu高性能算子的耗时比较：
已达到torch_npu高性能算子的70%左右的性能。

```
Kernel Output Match!
tilelang time: 0.8647611737251282 ms
torch time: 0.583781361579895 ms
```

如何调试TileLang Ascend算子

当前tilelang调试还需依赖AscendC调试工具来进行

步骤一：通过kernel.get_kernel_source()获取codegen后的AscendC源码

```
func = matmul(M, N, K, 128, 256, 64, 256, 2, 2)
print(func.get_kernel_source())
```

步骤二：编写静态调用脚本，在kernel入参中加入调试所需的workspace

```
dump_workspace = torch.empty(75 * 1024 * 1024 // 4).float().npu() # 初始化workspace空间
stream = torch.npu.current_stream()._as_parameter_

def tl_gemm():
    return lib.call(
        ctypes.c_void_p(a.data_ptr()),
        ctypes.c_void_p(b.data_ptr()),
        ctypes.c_void_p(c.data_ptr()),
        ctypes.c_void_p(dump_workspace.data_ptr()), # 加入到kernel入参中
        stream)

tl_gemm()
torch.npu.synchronize()
dump_workspace.detach().cpu().numpy().tofile("dump_workspace.bin") # 将结果作为.bin文件保存
```

如何调试TileLang Ascend算子

步骤三：将翻译后的AscendC代码加入相关初始化代码

```
extern "C" void call(uint8_t* A_handle, uint8_t* B_handle, uint8_t* C_handle, GM_ADDR dbgspace, actStream stream) {  
  
main_kernel<<<256, nullptr, stream>>>(A_handle, B_handle, C_handle, dbgspace, fftsAddr);  
  
#kernel 内部  
void main_kernel( GM_ADDR A_handle, GM_ADDR B_handle, GM_ADDR C_handle, GM_ADDR dbgspace, uint64_t fftsAddr) {  
    AscendC::InitDump(true, dump_workspace, 1024 * 1024);  
    AscendC::PRINTF("cid = %d\n", cid);  
}
```

步骤四：编译->执行->查看debug结果

```
bash build.sh example_gemm.cpp  
python test_example_gemm.py  
show_kernel_debug_data dump_workspace.bin /  
  
===== block.0 begin =====  
[Meta Info] block num: 256, core type: AIC, isMix: True  
cid = 206  
===== block.0 end =====  
===== block.1 begin =====  
[Meta Info] block num: 256, core type: AIC, isMix: True  
cid = 207  
===== block.1 end =====  
===== block.2 begin =====  
[Meta Info] block num: 256, core type: AIC, isMix: True  
cid = 208
```

参考用例：

https://github.com/tile-ai/tilelang-ascend/tree/ascendc_pto/examples/gemm_aot

目录

Part 1 TileLang介绍

Part 2 TileLang Ascend介绍

Part 3 TileLang Ascend开发入门

Part 4 TileLang Ascend 后续计划

TileLang Ascend后续计划

Framework Enhancement

- Automatic separation of cube and vector code, and elimination of explicit allocation of workspace on global memory
- Automatic pipeline injection and insertion of necessary synchronization between different hardware resources
- Automatic buffer reuse
- Dynamic Shape Support, including automatic TilingData extraction
- More complete tilelang primitives support
- PTO instruction support

Kernel Support

- Implementation of TileOPs using tilelang-ascend

Performance Optimization

- Optimization of instruction templates
- Complete integration into the existing network infrastructure
- Target comparable performance of corresponding AscendC programs

Tool development

- Profiling Tools for performance analysis

<https://gitcode.com/cann>



欢迎大家关注CANN开源社区和TileLang-Ascend开源社区！



<https://gitcode.com/cann/cann-recipes-infer/tree/master/ops/tilelang>



TileLang-Ascend
<https://github.com/tile-ai/tilelang-ascend>

Thank you.

社区愿景：打造开放易用、技术领先的AI算力新生态

社区使命：使能开发者基于CANN社区自主研究创新，构筑根深叶茂、跨产业协同共享共赢的CANN生态

Vision: Building an Open, Easy-to-Use, and Technology-leading AI Computing Ecosystem

Mission: Enable developers to independently research and innovate based on the CANN community and build a win-win CANN ecosystem with deep roots and cross-industry collaboration and sharing.



上CANN社区获取干货



关注CANN公众号获取资讯